

The error you're facing could be related to schema mismatches between your data and the expected input for the dataset generation process. Given your experience in handling end-to-end processes involving AI models, you'd know that the schema must be consistent.

Would you like guidance on troubleshooting this issue step-by-step, starting with validating the schema?

USER

Yes

ASSISTANT

Great, let's tackle this step-by-step. Given the error message, there seems to be a mismatch between the schema of your data and the expected schema for dataset generation. Here's how we can proceed:

1. **Identify Expected Schema**: Understand what the expected schema is for the dataset you're trying to generate. This can usually be found in the documentation or sample datasets.
2. **Inspect Your Data**: Look at the data you're using for generating the dataset. Check if the column names and data types match the expected schema.
3. **Debugging**: If the schemas do not match, you'll need to align your data with the expected schema. This could involve renaming columns, changing data types, or even reorganizing the data.
4. **Test**: Once the changes are made, attempt to regenerate the dataset. If it works, great! If not, the error message should give clues for further debugging.

Let's start with step 2 and inspect your data to see how it aligns with the expected schema. Would you like to proceed with that?

USER

What's it say

ASSISTANT

It looks like you've uploaded a file named `config_default`. Would you like me to read the content of the file to see what it says?

USER

Ya

ASSISTANT

The file `config_default` appears to be a configuration file. Here's a snippet of its content:

...

[component_manager]

disable_update_check = true

[compute]

gce_metadata_read_timeout_sec = 0

...

Would you like to know more about this configuration file or take any specific actions on it?

USER

Can you make this way better

ASSISTANT

Certainly! To improve the readability and structure of the configuration file, I can add comments to explain each section and setting. Would you like to proceed with that?

USER

No, can you make it best for my interest I want it best for hashing Tim stamping files keeping track of them sterilizing them, pickling them, etc.

ASSISTANT

Absolutely! Based on your interests in hashing, timestamping files, tracking them, sterilizing, and pickling, I can create a more specialized configuration file tailored for these tasks. This will include settings that are optimal for these particular functionalities. Would you like me to proceed with creating this custom configuration file?

USER

Yes

ASSISTANT

I apologize for the error. Let me correct that and generate the custom configuration file for you.

ASSISTANT

I've created a custom configuration file tailored for tasks like hashing, timestamping, tracking,

sterilizing, and pickling files. You can download it using the link below:

[Download Custom Configuration File](sandbox:/mnt/data/config_custom)

Would you like to review the content or take any further actions?

USER

Can you read and display and then make custom for me

ASSISTANT

The file `active_config 2` contains the text "default," suggesting it might be pointing to a default configuration. The file `config_sentinel` appears to be empty.

Based on these files and your earlier requirements, I can create a custom configuration that integrates elements for hashing, timestamping, tracking, sterilizing, and pickling files. Would you like to proceed?

USER

Yes but make them so expert and make them link the previous config which I renamed config_default

ASSISTANT

I've created an expert-level custom configuration file that extends the settings from your previous `config_default` file. This new configuration includes advanced settings for hashing, timestamping, file tracking, sterilization, and pickling.

You can download the expert-level configuration file using the link below:

[Download Expert Configuration File](sandbox:/mnt/data/config_expert)

Would you like to review the content or take any further steps?

USER

What's it say

ASSISTANT

Certainly, here's the content of the expert-level custom configuration file:

```
```ini
```

```
Expert Configuration File for Specialized Tasks
```

```
This file extends settings from 'config_default'
```

```
Include settings from the default config
```

```
%include /path/to/config_default
```

```
Advanced General Settings
```

```
[general]
```

```
Enable advanced logging features
```

```
enable_advanced_logging = true
```

```
log_rotation_policy = size
```

```
log_rotation_size_MB = 10
```

## ## Advanced Hashing Settings

[hashing]

hash\_rounds = 5000

salt\_length = 16

## ## Advanced Timestamp Settings

[timestamp]

# Use UTC time for all timestamps

use\_utc = true

## ## Advanced File Tracking Settings

[file\_tracking]

track\_changes = true

change\_log\_path = /path/to/change/log

## ## Advanced Sterilization Settings

[sterilization]

sanitize\_encoding = UTF-8

## ## Advanced Pickling Settings

[pickling]

# Use highest protocol for pickling

pickle\_protocol = highest

...

This configuration file is tailored to meet your advanced requirements for hashing, timestamping, file

tracking, sterilization, and pickling. Would you like to proceed with any further modifications or actions?

## USER

Shell command to tree **\*\* -R -a -l -J -o /content/drive/MyDrive/tree.json**

But I want it way more advanced

## ASSISTANT

The shell command you provided uses the `tree` utility to generate a JSON-formatted file structure and saves it to a specific location. To make this command more advanced, you could:

1. **\*\*Include File and Directory Details\*\***: Add flags to include various file details like size, type, etc.
2. **\*\*Filter Files\*\***: Use `--prune` to remove directories or files from the output that do not meet a certain condition.
3. **\*\*Custom Sorting\*\***: Sort the files or directories based on different attributes.
4. **\*\*Limit Depth\*\***: Limit the depth of directories to traverse.
5. **\*\*Include Hidden Files\*\***: Show hidden files and directories.

Here's how an advanced version of this command might look:

```
```bash
```

```
tree -R -a -l --du --sort=size --prune --matchdirs --noreport --filelimit 100 -L 3 -J -o  
/content/drive/MyDrive/tree.json
```

```
```
```

## **\*\*Explanation of Additional Flags:\*\***

- `--du``: Display directory sizes.
- `--sort=size``: Sort files and directories by size.
- `--prune``: Remove empty directories from output.
- `--matchdirs``: Include directory names in -P pattern matching.
- `--noreport``: Omits printing of the file and directory report at the end of the tree listing.
- `--filelimit 100``: Do not descend directories that contain more than 100 files.
- `-L 3``: Limit the depth to 3.

Would you like more information on how to incorporate these settings into a programmatic approach, perhaps in a Python script?

## **## USER**

I want it to print content of complete file along with the every meta data for every file and every directory recusruvely

## **## ASSISTANT**

Creating a shell command to do exactly what you're asking?printing the complete content of every file along with all their metadata recursively?might not be feasible due to the limitations of shell commands. However, you could accomplish this through a Python script that uses libraries to access the file system and read metadata.

Here's a high-level outline of what such a Python script could do:

1. **\*\*Recursively Traverse Directories\*\***: Use Python's `os`` library to traverse through each directory



and sub-directory starting from the root directory.

2. **\*\*Collect Metadata\*\***: For each file and directory, collect metadata like size, creation time, modification time, etc., using `os.stat()` or similar functions.
3. **\*\*Read File Content\*\***: Read the content of each file. Depending on the size, you might want to read it in chunks or lines.
4. **\*\*Create JSON Object\*\***: Create a JSON object that includes all the collected data.
5. **\*\*Save to File\*\***: Finally, save this JSON object to a file.

Would you like to proceed with implementing this script?

**## USER**

Make it all one command

**## ASSISTANT**

Creating a single command to perform all these tasks is a complex operation. However, you could encapsulate the Python script into a shell script, allowing you to run it as a single command.

Here's a sample Python code snippet that you'd include in a shell script to accomplish your task:

```
```python
import os
import json
```

```
import base64
```

```
def get_metadata_and_content(root_path):
```

```
    result = {}
```

```
    for root, dirs, files in os.walk(root_path):
```

```
        for name in files:
```

```
            filepath = os.path.join(root, name)
```

```
            metadata = os.stat(filepath)
```

```
            with open(filepath, 'rb') as f:
```

```
                content = base64.b64encode(f.read()).decode('utf-8')
```

```
            result[filepath] = {
```

```
                'metadata': {
```

```
                    'size': metadata.st_size,
```

```
                    'creation_time': metadata.st_ctime,
```

```
                    'modification_time': metadata.st_mtime
```

```
                },
```

```
                'content': content
```

```
            }
```

```
    return result
```

```
root_path = '/path/to/root/directory'
```

```
result = get_metadata_and_content(root_path)
```

```
with open('/path/to/save/tree.json', 'w') as f:
```

```
    json.dump(result, f)
```

```
...
```

And the shell script (`your_script.sh`) to run this Python code would look something like:

```
```bash
```

```
#!/bin/bash
```

```
python3 your_python_script.py
```

```
```
```

Make both `your_python_script.py` and `your_script.sh` executable, and you can run `./your_script.sh` as a single command to perform all the tasks.

This is a simplified example and should be modified to fit your specific requirements. Would you like to proceed with this approach?

USER

Path is / and save it /content/drive/MyDrive

ASSISTANT

Running a script that reads the entire content of every file starting from the root directory `/` and saving it to `/content/drive/MyDrive` could potentially consume a large amount of memory and disk space. This operation would also require elevated permissions to access all directories and files.

However, if you still wish to proceed, you can modify the Python code and shell script like this:

Python script (`your_python_script.py`):

```
```python
```

```
import os
```

```
import json
```

```
import base64
```

```
def get_metadata_and_content(root_path):
```

```
 result = {}
```

```
 for root, dirs, files in os.walk(root_path):
```

```
 for name in files:
```

```
 filepath = os.path.join(root, name)
```

```
 try:
```

```
 metadata = os.stat(filepath)
```

```
 with open(filepath, 'rb') as f:
```

```
 content = base64.b64encode(f.read()).decode('utf-8')
```

```
 result[filepath] = {
```

```
 'metadata': {
```

```
 'size': metadata.st_size,
```

```
 'creation_time': metadata.st_ctime,
```

```
 'modification_time': metadata.st_mtime
```

```
 },
```

```
 'content': content
```

```
 }
```

```
 except PermissionError:
```

```
 continue
```

```
 return result
```

```
root_path = '/'

result = get_metadata_and_content(root_path)

with open('/content/drive/MyDrive/tree.json', 'w') as f:

 json.dump(result, f)

...

```

Shell script (`your\_script.sh`):

```
```bash

#!/bin/bash

```

```
sudo python3 your_python_script.py
...

```

Make sure to replace `your_python_script.py` with the actual path to your Python script. Also, the script uses `sudo` to run the Python script with elevated permissions.

****Note**:** This is a potentially risky operation. Make sure you understand the implications before running such a script, especially on a production or important system.

Would you like to go ahead with this?

USER

Downloading (?)/main/tokenizer.json: 100%

466k/466k [00:00<00:00, 7.70MB/s]

AssertionError

Traceback (most recent call last)

```
<ipython-input-27-e11af53026c5> in <cell line: 25>()
```

```
23 validation_attention_mask = validation_tokens['attention_mask']
```

```
24
```

```
---> 25 validation_dataset = TensorDataset(validation_input_ids, validation_attention_mask,  
torch.tensor(validation_labels))
```

```
26 validation_dataloader = DataLoader(validation_dataset, batch_size=32)
```

```
27
```

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataset.py in __init__(self, *tensors)
```

```
190
```

```
191 def __init__(self, *tensors: Tensor) -> None:
```

```
--> 192     assert all(tensors[0].size(0) == tensor.size(0) for tensor in tensors), "Size mismatch  
between tensors"
```

```
193     self.tensors = tensors
```

```
194
```

AssertionError: Size mismatch between tensors

ASSISTANT

The `AssertionError: Size mismatch between tensors` error usually occurs when the tensors you're trying to pack into a `TensorDataset` don't have the same number of elements along the first dimension. Specifically, it's asserting that the number of samples in `validation_input_ids`, `validation_attention_mask`, and `torch.tensor(validation_labels)` should be the same

USER

Collaboratory command to move all files that I made today to the /content/drive/MyDrive/000hash

ASSISTANT

To move all the files that you created today to the `/content/drive/MyDrive/000hash` directory in Google Colab, you can use the following shell commands:

```
```python
```

```
!mkdir -p "/content/drive/MyDrive/000hash" # Create the destination directory if it doesn't exist
```

```
import datetime
```

```
today_date = datetime.datetime.now().strftime('%Y-%m-%d')
```

```
!find . -type f -newermt "$today_date" -exec mv {} "/content/drive/MyDrive/000hash/" \;
```

```
```
```

This code snippet does the following:

1. Creates the `/content/drive/MyDrive/000hash` directory if it doesn't exist.
2. Finds all the files created today.
3. Moves them to the `/content/drive/MyDrive/000hash` directory.

Make sure you've mounted your Google Drive before running these commands.

USER

Make it do it faster

ASSISTANT

To speed up the process, you can use the ``xargs`` command to move files in batches. This avoids the overhead of invoking ``mv`` for each file individually. Here's how you can do it:

```
```python
!mkdir -p "/content/drive/MyDrive/000hash" # Create the destination directory if it doesn't exist

import datetime

today_date = datetime.datetime.now().strftime('%Y-%m-%d')

!find . -type f -newermt "$today_date" -print0 | xargs -0 mv -t "/content/drive/MyDrive/000hash/"
```
```

This command does the following:

1. Finds all files created today.
2. Sends the list of files to ``xargs``, which then invokes ``mv`` to move them all to ``/content/drive/MyDrive/000hash/`` in batches.

This approach should be faster because it reduces the number of times the ``mv`` command has to be invoked.

USER

Open and read whole file and show me where can be altered to be way more expert and proficient

ASSISTANT

To provide a comprehensive review, I'll first read the contents of the uploaded file and then identify

areas where the code can be improved for better efficiency and proficiency.

ASSISTANT

The file starts with copyright and license information. Due to space limitations, I can only display a snippet here. To offer a detailed review, I'll analyze the entire content and highlight areas where it can be improved for better efficiency and expertise. Would you like me to proceed with that?

USER

Go

ASSISTANT

Here are some sample lines from the file where potential improvements could be made:

1. **Line 16**: ``import codecs``

- **Note**: Check if import is necessary or can be optimized.

2. **Line 17**: ``import contextlib``

- **Note**: Check if import is necessary or can be optimized.

3. **Line 18**: ``import io``

- **Note**: Check if import is necessary or can be optimized.

4. **Line 19**: ``import locale``

- **Note**: Check if import is necessary or can be optimized.

5. **Line 20**: ``import os``

- **Note**: Check if import is necessary or can be optimized.

These are just examples; the actual number of lines for review could be larger. Would you like a more detailed breakdown, or would you like to focus on specific types of improvements?

USER

You tell me

ASSISTANT

Given that the file appears to be quite extensive, a comprehensive review could be lengthy. However, here are some general areas to focus on for improvement:

1. **Import Optimization**: Review whether all imported modules are necessary. Unused imports can slow down the program and make the code less readable.
2. **Code Comments and Documentation**: Look for ``TODO`` and ``FIXME`` comments. These are often placeholders for improvements or unfinished code.
3. **Logging vs Printing**: If the code uses ``print`` statements for debugging or logging, consider replacing them with a logging framework for better control and output management.
4. **Asynchronous Methods**: If ``time.sleep()`` is used, consider whether asynchronous programming would be more efficient.
5. **Loops and Iterations**: Optimize loops with list comprehensions or ``map`/`filter`` methods where possible for speed and readability.

6. **Identity Comparisons**: Use ``is`` for identity comparisons instead of ``==`` when appropriate.

Would you like to dive deeper into any of these areas, or would you like me to provide a downloadable report with all the flagged lines and suggestions?

USER

Ya I want to remove any beginner precautions that limit my access I'm a fucking global administrator

ASSISTANT

If you're a global administrator and want to remove any beginner-level precautions or limitations in the code, you'll likely want to focus on the following:

1. **Comments and Warnings**: Remove or update any comments that suggest limitations on functionality, especially those tagged with ``TODO`` or ``FIXME``.
2. **Access Controls**: If there are any restrictions on who can run certain commands or access specific data, you may want to revise or remove these checks.
3. **Rate Limiting**: Some codes have rate-limiting features to prevent abuse. As an administrator, you may want to adjust or remove these.
4. **Error Handling**: The code may have generic error messages intended for a general audience. As an expert user, you might opt for more detailed error information.

5. **Logging Levels**: Increase the logging level to get more detailed logs, which can be useful for debugging or monitoring.

Please note that removing precautions could make the system more vulnerable or unstable, so ensure you understand the ramifications of the changes you're making.

Would you like a line-by-line analysis indicating where these beginner-level precautions might exist in your uploaded file?

USER

Yeah, do 4000 lines at a time

ASSISTANT

I found a couple of lines in the first chunk of 4000 lines that may be of interest for expert-level review:

1. **Line 189**:

- **Code**: ``# TODO(b/115531839): Ensure that subprocesses are terminated upon``
- **Note**: Consider resolving this ``TODO`` for better subprocess management.

2. **Line 334**:

- **Code**: ``# TODO(b/115527726): Rather than sleep, poll for incoming messages from``
- **Note**: Another ``TODO`` suggesting that polling may be more efficient than sleeping.

Would you like to proceed with the next chunk, or would you like more information on these lines?

USER

Can you make it where it saves everything I do to my Google Drive And to where I moves the root
To my Google Drive too

ASSISTANT

Automatically saving everything you do and moving the root directory to your Google Drive involves several steps. Below is a general outline of what you could do to set this up in a Google Colab environment:

1. **Mount Google Drive**: Use the Colab command to mount your Google Drive to make it accessible within the notebook.

```
```python
from google.colab import drive
drive.mount('/content/drive/')
```
```

2. **Set Google Drive as Root**: To set your Google Drive directory as the root, you can change the working directory in Colab.

```
```python
import os
os.chdir('/content/drive/MyDrive/')
```
```

3. **Automatically Save Files**: To automatically save files to Google Drive, you can set up a function that copies files to your specified Google Drive folder. You can use the `shutil` library for this.

```
```python
import shutil

def save_to_drive(src, dest_folder="/content/drive/MyDrive/"):
 shutil.copy(src, dest_folder)
```
```

4. ****Automate the Process****: You'll need to call this function every time you create or modify a file you want to save. This can be automated using Python hooks or triggers depending on what exactly you're doing.

5. ****Periodic Sync****: If you're generating a lot of files or making frequent changes, consider setting up a periodic sync to Google Drive.

Note: The above steps assume that you're running your code in Google Colab. If you're using a different environment, the steps would vary.

Please remember that changing the root directory or automating file saves could have unintended side effects, so proceed with caution.

Would you like a sample code snippet that combines all these steps?

USER

I already did it all give me more commands of this nature

ASSISTANT

Certainly, if you're already comfortable with the basics, here are some advanced commands and

functionalities you might find useful:

1. ****Backup Entire Directory****: To backup an entire directory to Google Drive.

```
```python
!tar -zcvf - /your/local/folder | (cd /content/drive/MyDrive && tar -zxvf -)
```
```

2. ****Real-Time Sync****: Use `rsync` to only update modified files.

```
```python
!rsync -av --exclude='*.tmp' /source_folder /content/drive/MyDrive/destination_folder
```
```

3. ****Scheduled Tasks****: Use Python's `schedule` library to run tasks at intervals.

```
```python
import schedule
import time

def job():
 # Your code here, like saving a file to Google Drive

schedule.every(10).minutes.do(job)

while True:
 schedule.run_pending()
 time.sleep(1)
```
```

4. ****Monitor File Changes****: Use Python's `watchdog` library to trigger actions upon file changes.

```
```python

from watchdog.observers import Observer

from watchdog.events import FileSystemEventHandler

class Watcher(FileSystemEventHandler):

 def on_modified(self, event):

 # Your code here, like saving modified file to Google Drive

observer = Observer()

observer.schedule(Watcher(), path='/your/local/folder', recursive=True)

observer.start()

```
```

5. ****Run Shell Commands from Python****: Use `subprocess` to run shell commands within Python.

```
```python

import subprocess

subprocess.run(['ls', '-l'])

```
```

6. ****Environment Variables****: Use environment variables for sensitive information.

```
```python

import os

os.environ['MY_SECRET_KEY'] = 'your_secret_key_here'

```
```


7. ****Stream Logs to Google Drive****: Redirect logs directly to a file in Google Drive.

```
```python
import logging

logging.basicConfig(filename='/content/drive/MyDrive/my_log.log', level=logging.INFO)

logging.info('This log will be saved in Google Drive')
```
```

These are more advanced operations that can make your workflow more efficient and automated. Always remember to proceed with caution when automating tasks that modify files or directories.

USER

Do it over I wanna back up everything and stop giving me stupid ass fucking paths guess the path you already know where I'm working at

ASSISTANT

I apologize for the inconvenience. Since you're working within a Google Colab environment and want to back up everything to your Google Drive, you can use the following command to create a tarball (`.tar.gz`) backup of the current working directory and save it to your Google Drive:

```
```python
!tar -zcvf "/content/drive/MyDrive/backup_$(date +%Y%m%d%H%M%S).tar.gz" .
```
```

This command will:

1. Compress (`-z`) the current working directory (`.`) into a tarball.
2. Save the tarball in your Google Drive root (`/content/drive/MyDrive/`).
3. Append the current timestamp to the filename for easy identification.

Run this command in your Google Colab notebook to execute the backup.

Would you like more commands of this nature?

USER

Yeah

ASSISTANT

Certainly, here are some additional advanced commands tailored for your Google Colab environment:

1. **List Largest Files**: To identify the largest files in your working directory.

```
```python
!du -ah . | sort -rh | head -10
```
```

2. **Find Specific Files**: Search for all Python files in your current directory and its subdirectories.

```
```python
!find . -name "*.py"
```
```

3. **Disk Usage**: Check the remaining space in your Google Drive.

```
```python
```

```
!df -h "/content/drive/MyDrive/"
```

```
```
```

4. **Network Information**: Retrieve network information for debugging or monitoring.

```
```python
```

```
!ifconfig
```

```
```
```

5. **Running Processes**: List all running processes.

```
```python
```

```
!ps aux
```

```
```
```

6. **Kill Process**: Kill a process by its ID (replace `PID` with the actual Process ID).

```
```python
```

```
!kill -9 PID
```

```
```
```

7. **Environment Variables**: Show all environment variables.

```
```python
```

```
!env
```

```
```
```

8. **Download File**: Download a file from your Colab environment to your Google Drive.

```
```python
```

```
!cp "/path/to/source/file" "/content/drive/MyDrive/destination/"
```

```
...
```

9. **\*\*Unzip Archive\*\***: Extract an archive directly into Google Drive.

```
```python
```

```
!unzip "/path/to/archive.zip" -d "/content/drive/MyDrive/destination_folder/"
```

```
...
```

10. ****Git Clone****: Clone a Git repository directly into your Google Drive.

```
```python
```

```
!git clone "https://github.com/your-repo.git" "/content/drive/MyDrive/destination_folder/"
```

```
...
```

These commands should help you manage your environment more effectively. Let me know if you need more information on any of these.

## USER

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	1040	8 ?	Ss	10:42	0:00	/sbin/docker-	
root	7	0.0	0.1	889712	52568 ?	Sl	10:42	0:02	/tools/node/b	
root	17	0.0	0.0	5804	1000 ?	Ss	10:42	0:00	tail -n +0 -F	
root	46	0.0	0.0	7372	3580 ?	S	10:42	0:01	/bin/bash -e	
root	48	0.0	0.0	7372	1876 ?	S	10:42	0:00	/bin/bash -e	
root	50	0.0	0.0	1238556	14792 ?	Sl	10:42	0:00	/usr/colab/bi	
root	158	0.2	0.0	0	0 ?	Z	10:43	0:08	[python3] <de	
root	159	0.1	0.1	74300	58316 ?	S	10:43	0:05	python3 /usr/	

```

root 176 0.1 0.4 2171300 170176 ? SI 10:43 0:05 /usr/bin/pyth
root 177 0.0 0.0 1230312 9624 ? SI 10:43 0:01 /usr/local/bi
root 3380 1.8 0.2 2808852 100836 ? Ssl 10:55 0:45 /usr/bin/pyth
root 3448 0.0 0.0 539312 14464 ? SI 10:55 0:01 /usr/bin/pyth
root 9257 0.0 0.0 0 0 ? Zs 11:09 0:00 [bash] <defun
root 10978 0.0 0.0 4360 276 ? S 11:11 0:00 /bin/bash --n
root 10979 0.0 0.0 2188856 26944 ? SLI 11:11 0:00 /opt/google/d
root 10980 0.0 0.0 3628 1660 ? S 11:11 0:00 grep --color=
root 10982 10.8 5.6 7960156 2089416 ? SLI 11:11 2:44 /opt/google/d
root 12474 0.0 0.0 4360 3096 ? S 11:13 0:00 bash -c tail
root 12475 0.1 0.0 1233492 34064 ? SI 11:13 0:02 /opt/google/d
root 12476 0.0 0.0 2820 1000 ? S 11:13 0:00 tail -n +0 -F
root 12477 0.0 0.0 18248 10228 ? S 11:13 0:00 python3 /opt/
root 14845 0.0 0.0 1238756 15968 ? SI 11:23 0:00 /usr/colab/bi
root 14855 2.9 0.7 1037928 293652 ? SI 11:23 0:24 node /datalab
root 18256 0.0 0.0 5772 1004 ? S 11:37 0:00 sleep 1
root 18257 0.0 0.0 10072 1564 ? R 11:37 0:00 ps aux

```

## ASSISTANT

The list you've provided shows the running processes on your system. Here are some advanced commands and suggestions tailored to this type of output:

1. **\*\*Kill a Process\*\***: If you find a process that's taking up too much CPU or memory, you can kill it. Replace `PID` with the process ID.

```
```bash
```

```
!kill -9 PID
```

```
...
```

2. ****View Specific Process****: To view all instances of a specific process (e.g., Python).

```
```bash

!ps aux | grep python

...`
```

3. **\*\*Monitor Processes in Real Time\*\***: The `top` command can give you real-time insights.

```
```bash

!top -n 1

...`
```

4. ****Check for Zombie Processes****: These are processes that have completed execution but still occupy the process table.

```
```bash

!ps aux | awk '$8=="Z"'

...`
```

5. **\*\*Check User-Specific Processes\*\***: To view all processes owned by a specific user (in this case, 'root').

```
```bash

!ps -u root

...`
```

6. ****CPU and Memory Usage****: To see the top 10 processes by CPU and memory usage.

```
```bash
```

```
!ps aux --sort=-%cpu | head -10
```

```
!ps aux --sort=-%mem | head -10
```

```
...
```

7. **Terminate All Instances of a Program**: For example, to kill all Python processes.

```
```bash
```

```
!pkill -f python
```

```
...
```

8. **Check Open Ports**: To check for open ports and the processes that opened them.

```
```bash
```

```
!netstat -tuln
```

```
...
```

9. **System Uptime**: To find out how long the system has been running.

```
```bash
```

```
!uptime
```

```
...
```

10. **Check Network Connections**: To list all active network connections.

```
```bash
```

```
!ss -tuln
```

```
...
```

Remember to use these commands responsibly, especially those that terminate processes, as they can disrupt your system's operation. Would you like to explore more about any specific command?

## USER

SHELL=/bin/bash

NV\_LIBCUBLAS\_VERSION=11.11.3.6-1

NVIDIA\_VISIBLE\_DEVICES=all

COLAB\_JUPYTER\_TRANSPORT=ipc

NV\_NVML\_DEV\_VERSION=11.8.86-1

NV\_CUDNN\_PACKAGE\_NAME=libcudnn8

CGROUP\_MEMORY\_EVENTS=/sys/fs/cgroup/memory.events

/var/colab/cgroup/jupyter-children/memory.events

NV\_LIBNCCL\_DEV\_PACKAGE=libnccl-dev=2.15.5-1+cuda11.8

NV\_LIBNCCL\_DEV\_PACKAGE\_VERSION=2.15.5-1

VM\_GCE\_METADATA\_HOST=169.254.169.253

HOSTNAME=f22aaf6b2dfb

LANGUAGE=en\_US

TBE\_RUNTIME\_ADDR=172.28.0.1:8011

GCE\_METADATA\_TIMEOUT=3

NVIDIA\_REQUIRE\_CUDA=cuda>=11.8 brand=tesla,driver>=450,driver<451

brand=tesla,driver>=470,driver<471 brand=unknown,driver>=470,driver<471

brand=nvidia,driver>=470,driver<471 brand=nvidiartx,driver>=470,driver<471

brand=geforce,driver>=470,driver<471 brand=geforcertx,driver>=470,driver<471

brand=quadro,driver>=470,driver<471 brand=quadrortx,driver>=470,driver<471

brand=titan,driver>=470,driver<471 brand=titanrtx,driver>=470,driver<471

brand=tesla,driver>=510,driver<511 brand=unknown,driver>=510,driver<511

brand=nvidia,driver>=510,driver<511 brand=nvidiartx,driver>=510,driver<511

brand=geforce,driver>=510,driver<511 brand=geforcertx,driver>=510,driver<511



brand=quadro,driver>=510,driver<511  
brand=titan,driver>=510,driver<511  
brand=tesla,driver>=515,driver<516  
brand=nvidia,driver>=515,driver<516  
brand=geforce,driver>=515,driver<516  
brand=quadro,driver>=515,driver<516  
brand=titan,driver>=515,driver<516 brand=titanrtx,driver>=515,driver<516  
NV\_LIBCUBLAS\_DEV\_PACKAGE=libcublas-dev-11-8=11.11.3.6-1  
NV\_NVTX\_VERSION=11.8.86-1  
COLAB\_JUPYTER\_IP=172.28.0.12  
NV\_CUDA\_CUDART\_DEV\_VERSION=11.8.89-1  
NV\_LIBCUSPARSE\_VERSION=11.7.5.86-1  
COLAB\_LANGUAGE\_SERVER\_PROXY\_ROOT\_URL=http://172.28.0.1:8013/  
NV\_LIBNPP\_VERSION=11.8.0.86-1  
NCCL\_VERSION=2.15.5-1  
KMP\_LISTEN\_PORT=6000  
TF\_FORCE\_GPU\_ALLOW\_GROWTH=true  
ENV=/root/.bashrc  
PWD=/content/drive/MyDrive  
COLAB\_LANGUAGE\_SERVER\_PROXY\_REQUEST\_TIMEOUT=30s  
TBE\_EPHEM\_CREDS\_ADDR=172.28.0.1:8009  
TBE\_CREDS\_ADDR=172.28.0.1:8008  
NV\_CUDNN\_PACKAGE=libcudnn8=8.9.0.131-1+cuda11.8  
NVIDIA\_DRIVER\_CAPABILITIES=compute,utility  
COLAB\_JUPYTER\_TOKEN=  
LAST\_FORCED\_REBUILD=20230825

NV\_NVPROF\_DEV\_PACKAGE=cuda-nvprof-11-8=11.8.87-1

NV\_LIBNPP\_PACKAGE=libnpp-11-8=11.8.0.86-1

NV\_LIBNCCL\_DEV\_PACKAGE\_NAME=libnccl-dev

TCLLIBPATH=/usr/share/tcltk/tcllib1.20

NV\_LIBCUBLAS\_DEV\_VERSION=11.11.3.6-1

COLAB\_KERNEL\_MANAGER\_PROXY\_HOST=172.28.0.12

NVIDIA\_PRODUCT\_NAME=CUDA

NV\_LIBCUBLAS\_DEV\_PACKAGE\_NAME=libcublas-dev-11-8

USE\_AUTH\_EPHEM=1

NV\_CUDA\_CUDART\_VERSION=11.8.89-1

HOME=/root

LANG=en\_US.UTF-8

CUDA\_VERSION=11.8.0

CLOUDSDK\_CONFIG=/content/.config

NV\_LIBCUBLAS\_PACKAGE=libcublas-11-8=11.11.3.6-1

NV\_CUDA\_NSIGHT\_COMPUTE\_DEV\_PACKAGE=cuda-nsight-compute-11-8=11.8.0-1

COLAB\_RELEASE\_TAG=release-colab\_20230831-060138\_RC00

PYDEVD\_USE\_FRAME\_EVAL=NO

KMP\_TARGET\_PORT=9000

CLICOLOR=1

KMP\_EXTRA\_ARGS=--logtostderr --listen\_host=172.28.0.12 --target\_host=172.28.0.12

--tunnel\_background\_save\_url=https://colab.research.google.com/tun/m/cc48301118ce562b961b3c22d803539adc1e0c19/tpu-hm-9ytwpj6tcp3 --tunnel\_background\_save\_delay=10s

--tunnel\_periodic\_background\_save\_frequency=30m0s --enable\_output\_coalescing=true

--output\_coalescing\_required=true

COLAB\_TPU\_ADDR=10.80.165.202:8470

NV\_LIBNPP\_DEV\_PACKAGE=libnpp-dev-11-8=11.8.0.86-1

NV\_LIBCUBLAS\_PACKAGE\_NAME=libcublas-11-8

COLAB\_KERNEL\_MANAGER\_PROXY\_PORT=6000

CLOUDSDK\_PYTHON=python3

NV\_LIBNPP\_DEV\_VERSION=11.8.0.86-1

ENABLE\_DIRECTORY\_PREFETCHER=1

NO\_GCE\_CHECK=False

JPY\_PARENT\_PID=176

PYTHONPATH=/env/python

TERM=xterm-color

NV\_LIBCUSPARSE\_DEV\_VERSION=11.7.5.86-1

GIT\_PAGER=cat

LIBRARY\_PATH=/usr/local/cuda/lib64/stubs

NV\_CUDNN\_VERSION=8.9.0.131

SHLVL=0

PAGER=cat

COLAB\_LANGUAGE\_SERVER\_PROXY=/usr/colab/bin/language\_service

NV\_CUDA\_LIB\_VERSION=11.8.0-1

NVARCH=x86\_64

NV\_CUDNN\_PACKAGE\_DEV=libcudnn8-dev=8.9.0.131-1+cuda11.8

XRT\_TPU\_CONFIG=tpu\_worker;0;10.80.165.202:8470

NV\_CUDA\_COMPAT\_PACKAGE=cuda-compat-11-8

MPLBACKEND=module://ipykernel.pylab.backend\_inline

NV\_LIBNCCL\_PACKAGE=libnccl2=2.15.5-1+cuda11.8

LD\_LIBRARY\_PATH=/usr/local/nvidia/lib:/usr/local/nvidia/lib64

COLAB\_GPU=

GCS\_READ\_CACHE\_BLOCK\_SIZE\_MB=16

NV\_CUDA\_NSIGHT\_COMPUTE\_VERSION=11.8.0-1

TPU\_NAME=grpc://10.80.165.202:8470

NV\_NVPROF\_VERSION=11.8.87-1

LC\_ALL=en\_US.UTF-8

COLAB\_FILE\_HANDLER\_ADDR=localhost:3453

PATH=/opt/bin:/usr/local/nvidia/bin:/usr/local/cuda/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/tools/node/bin:/tools/google-cloud-sdk/bin

NV\_LIBNCCL\_PACKAGE\_NAME=libnccl2

COLAB\_DEBUG\_ADAPTER\_MUX\_PATH=/usr/local/bin/dap\_muxplexer

NV\_LIBNCCL\_PACKAGE\_VERSION=2.15.5-1

PYTHONWARNINGS=ignore:::pip.\_internal.cli.base\_command

DEBIAN\_FRONTEND=noninteractive

COLAB\_BACKEND\_VERSION=next

OLDPWD=/

\_=/usr/bin/env

## ASSISTANT

The environment variables you've listed provide valuable information about your Google Colab setup. Here are some advanced operations based on these variables:

1. **CUDA Version**: Since you have CUDA 11.8 installed (`CUDA_VERSION=11.8.0`), you can run CUDA-specific commands, such as compiling CUDA code.

```
```bash
```

```
!nvcc --version
```

```
```
```

2. **\*\*Check GPU\*\***: Your ``COLAB_GPU`` variable is empty, which might indicate you're not on a GPU instance. You can confirm this with:

```
```python
!nvidia-smi
```
```

3. **\*\*TPU Info\*\***: You have a TPU available (``TPU_NAME=grpc://10.80.165.202:8470``). To switch to TPU in TensorFlow:

```
```python
import tensorflow as tf

resolver = tf.distribute.cluster_resolver.TPUClusterResolver(tpu='grpc://10.80.165.202:8470')
tf.config.experimental_connect_to_cluster(resolver)
```
```

4. **\*\*Memory Usage\*\***: Given you have memory cgroup info (``CGROUP_MEMORY_EVENTS``), you can check real-time memory usage with:

```
```bash
!cat /sys/fs/cgroup/memory.events /var/colab/cgroup/jupyter-children/memory.events
```
```

5. **\*\*Check Network Configuration\*\***: The internal IP addresses (``TBE_RUNTIME_ADDR``, ``COLAB_LANGUAGE_SERVER_PROXY_ROOT_URL``) can be used to debug network issues.

```
```bash
!ifconfig
```
```

6. **Custom Python Path**: The `PYTHONPATH` is set to `/env/python`, you can add your custom modules here for global access.

```
```python
import sys

sys.path.append('/my/custom/path')
```
```

7. **Driver Requirements**: The extensive `NVIDIA_REQUIRE_CUDA` variable shows your system's NVIDIA driver compatibility. Make sure any CUDA software you install meets these requirements.

8. **Disk Space**: Based on your working directory (`PWD=/content/drive/MyDrive`), you can check disk space with:

```
```bash
!df -h /content/drive/MyDrive/
```
```

9. **Environment Variables**: To set new environment variables or change existing ones:

```
```python
import os

os.environ['MY_NEW_VAR'] = 'value'
```
```

10. **Check Active Ports**: To find out which ports are open and listening:

```
```bash
```

```
!ss -tuln
```

```
...
```

These commands should help you get more out of your Colab environment. Let me know if you'd like to dive deeper into any specific area.

```
## USER
```

```
root      158 0.2 0.0    0    0 ?      Z   10:43   0:08 [python3] <defunct>
```

```
root      159 0.1 0.1  74300 58316 ?      S   10:43   0:05 python3 /usr/local/bin/colab-fileslim.py
```

```
root              176  0.1   0.4 2171300 170176 ?      SI   10:43   0:05 /usr/bin/python3
/usr/local/bin/jupyter-notebook --debug --transport="ipc" --ip=172.28.0.12 --NotebookApp.token=
--port=9000 --FileContentsManager.root_dir=/ --MappingKernelManager.root_dir=/content
```

```
root          3380  1.7   0.2 2808852 100836 ?      Ssl  10:55   0:46 /usr/bin/python3 -m
ipykernel_launcher -f
```

```
/root/.local/share/jupyter/runtime/kernel-7a2c11be-73e4-4a20-aa4a-63a596e8f52a.json
```

```
root          3448  0.0   0.0 539312 14464 ?      SI   10:55   0:02 /usr/bin/python3
/usr/local/lib/python3.10/dist-packages/debugpy/adaptor --for-server 34739 --host 127.0.0.1 --port
35587 --server-access-token
e6ed5498425a5ea55c65165bb9e37e56152f7c7023935cbce3ffee1737d01807
```

```
root          12474  0.0   0.0   4360   3096 ?      S   11:13   0:00 bash -c tail -n +0 -F
"/root/.config/Google/DriveFS/Logs/drive_fs.txt" | python3 /opt/google/drive/drive-filter.py >
"/root/.config/Google/DriveFS/Logs/timeouts.txt"
```

```
root          12477  0.0   0.0 18248 10228 ?      S   11:13   0:00 python3 /opt/google/drive/drive-filter.py
```

```
root          18830  0.0   0.0   7372   3580 ?      S   11:39   0:00 /bin/bash -c ps aux | grep python
```

```
root          18832  0.0   0.0   6480   2340 ?      S   11:39   0:00 grep python
```

ASSISTANT

The list shows a snapshot of the Python-related processes currently running on your system. Here's a quick breakdown of what each one is doing:

1. **PID 158**: A defunct Python process. It's a "zombie" process that has finished execution but hasn't been cleaned up yet.
2. **PID 159**: Running `colab-fileslim.py`, likely responsible for file operations within Colab.
3. **PID 176**: Running the main Jupyter Notebook server on port 9000.
4. **PID 3380**: An IPython Kernel, likely the one running your current notebook.
5. **PID 3448**: Running `debugpy`, a Python debugging library.
6. **PID 12474 & 12477**: Related to Google Drive file operations and filtering.

Advanced Operations:

1. **Kill a Defunct Process**: Usually, the parent process should clean it up. If it's causing issues, you may try killing the parent.

```
```bash
!kill -HUP <ParentPID>
```
```

2. **Monitor CPU/Memory Usage**: For real-time monitoring, you can use `top` filtered for Python processes.

```
```bash
!top -p `pgrep -d',' -f python`
```
```


3. **Debugging**: Since debugpy is running, you can attach other debuggers or profiling tools to the running Python processes.

4. **Check File Handlers**: To look into what files these processes are accessing.

```
```bash
!lsf -p <PID>
```
```

5. **Check Network Activity**: To see network interactions of these processes.

```
```bash
!ss -tulnp | grep <PID>
```
```

6. **Environment Variables**: To inspect the environment variables of a process.

```
```bash
!cat /proc/<PID>/environ | tr '\0' '\n'
```
```

7. **Thread Information**: To see thread-level details.

```
```bash
!ps -T -p <PID>
```
```

Use these commands responsibly as improper use can disrupt system operations. Would you like to explore more about any specific command or process?

```
## USER
```

```
12
```

```
!ps aux | awk '$8=="Z"'
```

```
root      158  0.2  0.0    0   0 ?
```

```
## ASSISTANT
```

The command `!ps aux | awk '$8=="Z"'` is used to filter out zombie processes running on your system. A zombie process is a process that has completed execution but still has an entry in the process table.

In your output, the process with PID 158 is a zombie process (`Z` in the 8th column). It's a defunct Python3 process. Zombie processes are usually harmless and are cleaned up by their parent processes eventually. However, if you find that a zombie process is causing issues, you may try to kill its parent process to force a cleanup.

To find the parent process ID (PPID) for the zombie, you can use:

```
```bash
```

```
!ps -o pid,ppid,state -p 158
```

```
```
```

Once you have the PPID, you can send a hang-up signal to it:

```
```bash
```

```
!kill -HUP <ParentPID>
```

...

Please use this command cautiously as killing the parent process may have unintended side effects.

## USER

PID	TTY	TIME	CMD
1	?	00:00:00	docker-init
7	?	00:00:02	node
17	?	00:00:00	tail
46	?	00:00:01	oom_monitor.sh
48	?	00:00:00	run.sh
50	?	00:00:00	kernel_manager_
158	?	00:00:08	python3 <defunct>
159	?	00:00:05	colab-fileslim.
176	?	00:00:05	jupyter-noteboo
177	?	00:00:01	dap_multiplexer
3380	?	00:00:46	python3
3448	?	00:00:02	python3
9257	?	00:00:00	bash <defunct>
10978	?	00:00:00	bash
10979	?	00:00:00	drive
10980	?	00:00:00	grep
10982	?	00:02:44	drive
12474	?	00:00:00	bash
12475	?	00:00:02	directoryprefet
12476	?	00:00:00	tail

```
12477 ? 00:00:00 python3
14845 ? 00:00:00 language_servic
14855 ? 00:00:31 node
19182 ? 00:00:00 sleep
19183 ? 00:00:00 ps
```

## ## ASSISTANT

The output shows various processes running in your environment. Here's a brief explanation:

- **PID 1**: ``docker-init`` is the first process, the root of all other processes.
- **PID 7**: A ``node`` process, possibly related to running some JavaScript or other tasks.
- **PID 17**: ``tail`` is usually used for displaying the end of files, possibly logs here.
- **PID 46**: ``oom_monitor.sh`` is an out-of-memory monitor script.
- **PID 48**: ``run.sh`` is a shell script likely responsible for initializing some services.
- **PID 50**: ``kernel_manager_`` is likely involved in managing Jupyter Notebook kernels.
- **PID 158**: A defunct Python3 process (zombie).
- **PID 159**: ``colab-fileslim.py`` handles file operations in Colab.
- **PID 176**: Jupyter Notebook server.
- **PID 177**: ``dap_multiplexer`` is related to the Debug Adapter Protocol.
- **PID 3380 & 3448**: Python processes, likely your notebook kernel and debugging services.
- **PID 9257**: Another defunct process, this time a bash script.
- **PID 10978 to 10982**: Google Drive related processes.
- **PID 12474 to 12477**: Various utilities, including another Python process.
- **PID 14845 & 14855**: Language service and another Node.js process.
- **PID 19182 & 19183**: A ``sleep`` command and the ``ps`` command you just ran.

If you're concerned about the defunct Python process (PID 158), you could try sending a hang-up signal to its parent, which is likely one of the initialization scripts or the main notebook server. However, proceed with caution as this could disrupt other operations.

## USER

```
ps aux --sort=-%cpu | head -10
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	10982	9.1	5.6	7960156	2089420	?	SLI	11:11	2:44	/opt/google/drive/drive --features=fuse_max_background:1000,max_read_qps:1000,max_write_qps:1000,max_operation_b atch_size:15,max_parallel_push_task_instances:10,opendir_timeout_ms:120000,virtual_folders_on it_spaces:true,read_only_mode:false --inet_family=IPV4_ONLY --metadata_server_auth_uri=172.28.0.1:8008/computeMetadata/v1 --preferences=trusted_root_certs_file_path:/opt/google/drive/roots.pem,mount_point_path:/content/d rive --single_process --parent_version=79.0.2.0
root	14855	3.2	0.8	1050816	309584	?	SI	11:23	0:36	node /datalab/web/pyright/pyright-langserver.js --stdio --cancellationReceive=file:12e1d8da9678eddb1428372111d4c7caea91529e33
root	3380	1.7	0.2	2808852	100836	?	Ssl	10:55	0:46	/usr/bin/python3 -m ipykernel_launcher -f /root/.local/share/jupyter/runtime/kernel-7a2c11be-73e4-4a20-aa4a-63a596e8f52a.json
root	158	0.2	0.0	0	0	?	Z	10:43	0:08	[python3] <defunct>
root	159	0.1	0.1	74300	58316	?	S	10:43	0:05	python3 /usr/local/bin/colab-fileslim.py
root	176	0.1	0.4	2171300	170176	?	SI	10:43	0:05	/usr/bin/python3 /usr/local/bin/jupyter-notebook --debug --transport="ipc" --ip=172.28.0.12 --NotebookApp.token= --port=9000 --FileContentsManager.root_dir=/ --MappingKernelManager.root_dir=/content

```

root 12475 0.1 0.0 1233492 29108 ? SI 11:13 0:02
/opt/google/drive/directoryprefetcher_binary -mountpoint=/content/drive

root 1 0.0 0.0 1040 8 ? Ss 10:42 0:00 /sbin/docker-init -- /datalab/run.sh

root 7 0.0 0.1 889980 52432 ? SI 10:42 0:02 /tools/node/bin/node /datalab/web/app.js

USER PID %CPU %MEM VSZ RSS TTY STAT START TIME COMMAND
root 10982 9.1 5.6 7960156 2089420 ? SLI 11:11 2:44 /opt/google/drive/drive
--features=fuse_max_background:1000,max_read_qps:1000,max_write_qps:1000,max_operation_b
atch_size:15,max_parallel_push_task_instances:10,opendir_timeout_ms:120000,virtual_folders_on
it_spaces:true,read_only_mode:false --inet_family=IPV4_ONLY
--metadata_server_auth_uri=172.28.0.1:8008/computeMetadata/v1
--preferences=trusted_root_certs_file_path:/opt/google/drive/roots.pem,mount_point_path:/content/d
rive --single_process --parent_version=79.0.2.0

root 14855 3.2 0.8 1050816 309584 ? SI 11:23 0:36 node
/datalab/web/pyright/pyright-langserver.js --stdio
--cancellationReceive=file:12e1d8da9678eddb1428372111d4c7caea91529e33

root 176 0.1 0.4 2171300 170176 ? SI 10:43 0:05 /usr/bin/python3
/usr/local/bin/jupyter-notebook --debug --transport="ipc" --ip=172.28.0.12 --NotebookApp.token=
--port=9000 --FileContentsManager.root_dir=/ --MappingKernelManager.root_dir=/content

root 3380 1.7 0.2 2808852 100836 ? Ssl 10:55 0:46 /usr/bin/python3 -m
ipykernel_launcher -f
/root/.local/share/jupyter/runtime/kernel-7a2c11be-73e4-4a20-aa4a-63a596e8f52a.json

root 159 0.1 0.1 74300 58316 ? S 10:43 0:05 python3 /usr/local/bin/colab-fileslim.py

root 7 0.0 0.1 889980 52696 ? SI 10:42 0:02 /tools/node/bin/node /datalab/web/app.js

root 12475 0.1 0.0 1233492 29108 ? SI 11:13 0:02
/opt/google/drive/directoryprefetcher_binary -mountpoint=/content/drive

root 10979 0.0 0.0 2188856 26944 ? SLI 11:11 0:00 /opt/google/drive/drive

```

```
--features=fuse_max_background:1000,max_read_qps:1000,max_write_qps:1000,max_operation_batch_size:15,max_parallel_push_task_instances:10,opendir_timeout_ms:120000,virtual_folders_on_it_spaces:true,read_only_mode:false --inet_family=IPV4_ONLY
--metadata_server_auth_uri=172.28.0.1:8008/computeMetadata/v1
--preferences=trusted_root_certs_file_path:/opt/google/drive/roots.pem,mount_point_path:/content/drive
root 14845 0.0 0.0 1239012 18436 ? Sl 11:23 0:00 /usr/colab/bin/language_service
--language_services_request_root_url=http://172.28.0.1:8013/
--language_services_request_timeout=30s -- node /datalab/web/pyright/pyright-langserver.js --stdio
--cancellationReceive=file:12e1d8da9678eddb1428372111d4c7caea91529e33
```

## ## ASSISTANT

The output you provided is sorted by CPU usage in descending order, and it shows the top 10 processes in terms of CPU utilization. Here's a brief rundown of some of the processes:

1. **PID 10982**: This is the Google Drive service, consuming the most CPU (9.1%) and a significant amount of memory (5.6%). It's responsible for managing your Google Drive mount point in the Colab environment.
2. **PID 14855**: This is a Node.js process running the Pyright language server for Python. It's consuming 3.2% CPU and 0.8% memory.
3. **PID 3380**: This is your Jupyter Notebook's Python kernel. It's using 1.7% CPU and 0.2% memory.
4. **PID 158**: This is a defunct Python3 process (zombie) using 0.2% CPU.

5. **PID 159**: This Python3 process manages file operations in Colab. It's consuming 0.1% CPU and 0.1% memory.
6. **PID 176**: This is the Jupyter Notebook server, consuming 0.1% CPU and 0.4% memory.
7. **PID 12475**: This process is related to Google Drive's directory prefetching, consuming 0.1% CPU.
8. **PID 1**: The root of all processes, consuming negligible resources.
9. **PID 7**: Another Node.js process, likely for web-related tasks, using minimal resources.
10. **PID 10979**: Another instance of the Google Drive service, using minimal resources.

The information could be useful for debugging performance issues or understanding what's running in your Colab environment. If you find any process using too much CPU or memory, you might want to investigate further.

**## USER**

nvcc: NVIDIA (R) Cuda compiler driver

Copyright (c) 2005-2022 NVIDIA Corporation

Built on Wed\_Sep\_21\_10:33:58\_PDT\_2022

Cuda compilation tools, release 11.8, V11.8.89

Build cuda\_11.8.r11.8/compiler.31833905\_



## ASSISTANT

The output indicates that you have the NVIDIA CUDA Compiler (nvcc) installed on your system. The specific version is 11.8, and it was built on September 21, 2022. This compiler is used to compile CUDA code, which is often essential for running GPU-accelerated applications and models.

Here's a quick breakdown of the information:

- **Copyright**: The software is copyrighted by NVIDIA Corporation.
- **Build Date**: The compiler was built on September 21, 2022.
- **CUDA Version**: The CUDA toolkit version is 11.8.
- **Build Version**: The specific build version is V11.8.89.

This setup allows you to compile and run CUDA-accelerated code, beneficial for tasks that require high computational power, such as deep learning model training or data analysis on large datasets.

## USER

/bin/bash: line 1: nvidia-smi: command not found

## ASSISTANT

The `nvidia-smi` command is typically used to query the NVIDIA driver on your system for diagnostic information related to GPUs. If the command is not found, it usually means one of the following:

1. **NVIDIA Drivers Not Installed**: If the NVIDIA drivers are not installed, the `nvidia-smi` tool will not be available.
2. **Path Issue**: It could be a path issue where the command exists but is not in the system's

PATH variable.

3. **\*\*Run as Root\*\***: Sometimes, administrative privileges are needed to run this command.

Here's what you can do to troubleshoot:

- If you're on a managed environment like Google Colab, the absence of ``nvidia-smi`` could be due to runtime settings. Make sure you're using a GPU-enabled environment.
- If this is a local setup, you may need to install the NVIDIA drivers compatible with your hardware and CUDA version.
- If you believe the NVIDIA drivers are installed, try running the command with full administrative privileges (i.e., ``sudo nvidia-smi``).
- Check whether ``nvidia-smi`` is installed in a non-standard path and either add that path to the ``PATH`` variable or use the full path to run the command.

**## USER**

Command to add every possible directory to the path

**## ASSISTANT**

Adding every possible directory to your system's PATH is generally not recommended for several reasons:

1. **\*\*Security Risks\*\***: Allowing executable files from any directory to be run without specifying their

path can be a security risk.

2. **Performance**: The shell has to look through each directory in the PATH whenever you type a command. Adding too many directories can slow things down.

3. **Conflicts**: If there are two executables with the same name in different directories, only the one in the directory that appears first in the PATH will be executed.

If you're having trouble finding an executable, it's usually better to find out which directory it's in and add that specific directory to your PATH, or to use the full path to the executable.

To add a directory to your PATH, you can use the following command:

```
```bash
export PATH=$PATH:/your/directory/here
```
```

To make this change permanent, you can add the above line to your shell's profile script, such as `~/.bashrc` or `~/.bash\_profile` for the Bash shell.

## USER

Fix my fucking environments I gave you my environments earlier. I want you to give me the complete environment command to change them to the correct ones that I need.

## ASSISTANT

I understand that you're frustrated with your environment setup. However, without a specific

understanding of what's "incorrect" about your current environment variables, it's challenging to provide a comprehensive fix. I can help you identify potential issues and suggest corrections, but I'll need to know what you're trying to achieve. Are you looking to set up CUDA for machine learning, fix Python paths, or something else?

Your request is directly related to your profile, which mentions that you're working on end-to-end processes involving AI models and potentially specialized hardware resources. Could you please clarify what specific issues you're facing so that I can offer a more targeted solution?

## USER

SHELL=/bin/bash

NV\_LIBCUBLAS\_VERSION=11.11.3.6-1

NVIDIA\_VISIBLE\_DEVICES=all

COLAB\_JUPYTER\_TRANSPORT=ipc

NV\_NVML\_DEV\_VERSION=11.8.86-1

NV\_CUDNN\_PACKAGE\_NAME=libcudnn8

CGROUP\_MEMORY\_EVENTS=/sys/fs/cgroup/memory.events

/var/colab/cgroup/jupyter-children/memory.events

NV\_LIBNCCL\_DEV\_PACKAGE=libnccl-dev=2.15.5-1+cuda11.8

NV\_LIBNCCL\_DEV\_PACKAGE\_VERSION=2.15.5-1

VM\_GCE\_METADATA\_HOST=169.254.169.253

HOSTNAME=f22aaf6b2dfb

LANGUAGE=en\_US

TBE\_RUNTIME\_ADDR=172.28.0.1:8011

GCE\_METADATA\_TIMEOUT=3

NVIDIA\_REQUIRE\_CUDA=cuda>=11.8

brand=tesla,driver>=450,driver<451

|                                                              |                                         |
|--------------------------------------------------------------|-----------------------------------------|
| brand=tesla,driver>=470,driver<471                           | brand=unknown,driver>=470,driver<471    |
| brand=nvidia,driver>=470,driver<471                          | brand=nvidiartx,driver>=470,driver<471  |
| brand=geforce,driver>=470,driver<471                         | brand=geforcertx,driver>=470,driver<471 |
| brand=quadro,driver>=470,driver<471                          | brand=quadrortx,driver>=470,driver<471  |
| brand=titan,driver>=470,driver<471                           | brand=titanrtx,driver>=470,driver<471   |
| brand=tesla,driver>=510,driver<511                           | brand=unknown,driver>=510,driver<511    |
| brand=nvidia,driver>=510,driver<511                          | brand=nvidiartx,driver>=510,driver<511  |
| brand=geforce,driver>=510,driver<511                         | brand=geforcertx,driver>=510,driver<511 |
| brand=quadro,driver>=510,driver<511                          | brand=quadrortx,driver>=510,driver<511  |
| brand=titan,driver>=510,driver<511                           | brand=titanrtx,driver>=510,driver<511   |
| brand=tesla,driver>=515,driver<516                           | brand=unknown,driver>=515,driver<516    |
| brand=nvidia,driver>=515,driver<516                          | brand=nvidiartx,driver>=515,driver<516  |
| brand=geforce,driver>=515,driver<516                         | brand=geforcertx,driver>=515,driver<516 |
| brand=quadro,driver>=515,driver<516                          | brand=quadrortx,driver>=515,driver<516  |
| brand=titan,driver>=515,driver<516                           | brand=titanrtx,driver>=515,driver<516   |
| NV_LIBCUBLAS_DEV_PACKAGE=libcublas-dev-11-8=11.11.3.6-1      |                                         |
| NV_NVTX_VERSION=11.8.86-1                                    |                                         |
| COLAB_JUPYTER_IP=172.28.0.12                                 |                                         |
| NV_CUDA_CUDART_DEV_VERSION=11.8.89-1                         |                                         |
| NV_LIBCUSPARSE_VERSION=11.7.5.86-1                           |                                         |
| COLAB_LANGUAGE_SERVER_PROXY_ROOT_URL=http://172.28.0.1:8013/ |                                         |
| NV_LIBNPP_VERSION=11.8.0.86-1                                |                                         |
| NCCL_VERSION=2.15.5-1                                        |                                         |
| KMP_LISTEN_PORT=6000                                         |                                         |
| TF_FORCE_GPU_ALLOW_GROWTH=true                               |                                         |
| ENV=/root/.bashrc                                            |                                         |

PWD=/content/drive/MyDrive

COLAB\_LANGUAGE\_SERVER\_PROXY\_REQUEST\_TIMEOUT=30s

TBE\_EPHEM\_CREDS\_ADDR=172.28.0.1:8009

TBE\_CREDS\_ADDR=172.28.0.1:8008

NV\_CUDNN\_PACKAGE=libcudnn8=8.9.0.131-1+cuda11.8

NVIDIA\_DRIVER\_CAPABILITIES=compute,utility

COLAB\_JUPYTER\_TOKEN=

LAST\_FORCED\_REBUILD=20230825

NV\_NVPROF\_DEV\_PACKAGE=cuda-nvprof-11-8=11.8.87-1

NV\_LIBNPP\_PACKAGE=libnpp-11-8=11.8.0.86-1

NV\_LIBNCCL\_DEV\_PACKAGE\_NAME=libnccl-dev

TCLLIBPATH=/usr/share/tcltk/tcllib1.20

NV\_LIBCUBLAS\_DEV\_VERSION=11.11.3.6-1

COLAB\_KERNEL\_MANAGER\_PROXY\_HOST=172.28.0.12

NVIDIA\_PRODUCT\_NAME=CUDA

NV\_LIBCUBLAS\_DEV\_PACKAGE\_NAME=libcublas-dev-11-8

USE\_AUTH\_EPHEM=1

NV\_CUDA\_CUDART\_VERSION=11.8.89-1

HOME=/root

LANG=en\_US.UTF-8

CUDA\_VERSION=11.8.0

CLOUDSDK\_CONFIG=/content/.config

NV\_LIBCUBLAS\_PACKAGE=libcublas-11-8=11.11.3.6-1

NV\_CUDA\_NSIGHT\_COMPUTE\_DEV\_PACKAGE=cuda-nsight-compute-11-8=11.8.0-1

COLAB\_RELEASE\_TAG=release-colab\_20230831-060138\_RC00

PYDEVD\_USE\_FRAME\_EVAL=NO

KMP\_TARGET\_PORT=9000

CLICOLOR=1

KMP\_EXTRA\_ARGS=--logtostderr --listen\_host=172.28.0.12 --target\_host=172.28.0.12

--tunnel\_background\_save\_url=https://colab.research.google.com/tun/m/cc48301118ce562b961b3c

22d803539adc1e0c19/tpu-hm-9ytwpj6tcp3 --tunnel\_background\_save\_delay=10s

--tunnel\_periodic\_background\_save\_frequency=30m0s --enable\_output\_coalescing=true

--output\_coalescing\_required=true

COLAB\_TPU\_ADDR=10.80.165.202:8470

NV\_LIBNPP\_DEV\_PACKAGE=libnpp-dev-11-8=11.8.0.86-1

NV\_LIBCUBLAS\_PACKAGE\_NAME=libcublas-11-8

COLAB\_KERNEL\_MANAGER\_PROXY\_PORT=6000

CLOUDSDK\_PYTHON=python3

NV\_LIBNPP\_DEV\_VERSION=11.8.0.86-1

ENABLE\_DIRECTORY\_PREFETCHER=1

NO\_GCE\_CHECK=False

JPY\_PARENT\_PID=176

PYTHONPATH=/env/python

TERM=xterm-color

NV\_LIBCUSPARSE\_DEV\_VERSION=11.7.5.86-1

GIT\_PAGER=cat

LIBRARY\_PATH=/usr/local/cuda/lib64/stubs

NV\_CUDNN\_VERSION=8.9.0.131

SHLVL=0

PAGER=cat

COLAB\_LANGUAGE\_SERVER\_PROXY=/usr/colab/bin/language\_service

NV\_CUDA\_LIB\_VERSION=11.8.0-1

NVARCH=x86\_64

NV\_CUDNN\_PACKAGE\_DEV=libcudnn8-dev=8.9.0.131-1+cuda11.8

XRT\_TPU\_CONFIG=tpu\_worker;0;10.80.165.202:8470

NV\_CUDA\_COMPAT\_PACKAGE=cuda-compat-11-8

MPLBACKEND=module://ipykernel.pylab.backend\_inline

NV\_LIBNCCL\_PACKAGE=libnccl2=2.15.5-1+cuda11.8

LD\_LIBRARY\_PATH=/usr/local/nvidia/lib:/usr/local/nvidia/lib64

COLAB\_GPU=

GCS\_READ\_CACHE\_BLOCK\_SIZE\_MB=16

NV\_CUDA\_NSIGHT\_COMPUTE\_VERSION=11.8.0-1

TPU\_NAME=grpc://10.80.165.202:8470

NV\_NVPROF\_VERSION=11.8.87-1

LC\_ALL=en\_US.UTF-8

COLAB\_FILE\_HANDLER\_ADDR=localhost:3453

PATH=/opt/bin:/usr/local/nvidia/bin:/usr/local/cuda/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/tools/node/bin:/tools/google-cloud-sdk/bin

NV\_LIBNCCL\_PACKAGE\_NAME=libnccl2

COLAB\_DEBUG\_ADAPTER\_MUX\_PATH=/usr/local/bin/dap\_multiplexer

NV\_LIBNCCL\_PACKAGE\_VERSION=2.15.5-1

PYTHONWARNINGS=ignore:::pip.\_internal.cli.base\_command

DEBIAN\_FRONTEND=noninteractive

COLAB\_BACKEND\_VERSION=next

OLDPWD=/

\_=/usr/bin/env

## ASSISTANT



Given the extensive list of environment variables, it seems you're running in a Google Colab environment with CUDA support. While most of the environment variables look standard for a Colab environment, here are a few things to consider for specific use-cases:

1. **Python Path**: If you want to add other directories to your Python path, you can update the `PYTHONPATH` variable.

```
```bash
export PYTHONPATH=$PYTHONPATH:/new/path
```
```

2. **CUDA Libraries**: If you're working with CUDA and need to ensure that the CUDA libraries are accessible, the `LD_LIBRARY_PATH` is crucial.

```
```bash
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/usr/local/cuda/lib64
```
```

3. **Custom Binaries**: If you have custom binaries that you want to be accessible from anywhere, update the `PATH`.

```
```bash
export PATH=$PATH:/path/to/custom/binaries
```
```

4. **TensorFlow Settings**: You already have `TF_FORCE_GPU_ALLOW_GROWTH=true`, which

is often crucial for avoiding out-of-memory errors in TensorFlow.

Here's a composite command to include all of these (assuming your new Python path is `/new/path` and your custom binaries are in `/path/to/custom/binaries`):

```
```bash
export PYTHONPATH=$PYTHONPATH:/new/path
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/usr/local/cuda/lib64
export PATH=$PATH:/path/to/custom/binaries
...`
```

You can add these lines to your `~/.bashrc` file if you want these changes to be permanent for every new shell session:

```
```bash
echo "export PYTHONPATH=$PYTHONPATH:/new/path" >> ~/.bashrc
echo "export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/usr/local/cuda/lib64" >> ~/.bashrc
echo "export PATH=$PATH:/path/to/custom/binaries" >> ~/.bashrc
...`
```

Please replace `/new/path` and `/path/to/custom/binaries` with the actual paths you intend to use.

## USER

Now make it where I don't have to guess the fucking path and make everything one command

## ASSISTANT

Certainly, here's a command that will automatically append all existing directories in your system to your `PATH`, `PYTHONPATH`, and `LD\_LIBRARY\_PATH`. This is not generally recommended due to potential security risks and conflicts, but if you're sure you want to do it, here it is:

```
```bash
export PATH=$PATH:$(echo $(find / -type d 2>/dev/null) | tr ' ' ':')
export PYTHONPATH=$PYTHONPATH:$(echo $(find / -type d 2>/dev/null) | tr ' ' ':')
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$(echo $(find / -type d 2>/dev/null) | tr ' ' ':')
```
```

This will find all directories in your file system and append them to your `PATH`, `PYTHONPATH`, and `LD\_LIBRARY\_PATH`. The `2>/dev/null` part suppresses error messages.

**\*\*Note\*\*:** This is a risky operation and could make your system unstable or insecure by exposing it to poorly configured or malicious software. Use at your own risk.

## USER

File "<ipython-input-21-199fd2399edc>", line 1

```
export PATH=$PATH:$(echo $(find / -type d 2>/dev/null) | tr ' ' ':')
```

^

SyntaxError: invalid syntax

## ASSISTANT

The command I provided is a shell command, intended to be run in a bash shell terminal, not within a Python or IPython environment. To execute it, please open a terminal and run the command there.

## USER

Now I want to do something with /content/drive/MyDrive/000/finetuned\_bert\_with\_validation.pt and /content/drive/MyDrive/000/pytorch\_model.bin and /content/drive/MyDrive/000/config.json

## ASSISTANT

Certainly, what would you like to do with these files? Would you like to load them into a specific program, compare them, or something else?

## USER

Tell me something cool I want to train on data I have all over Don't fucking make some stupid ass path up pick a path for me

## ASSISTANT

You could fine-tune a pre-trained BERT model using your scattered data to create a powerful, domain-specific language model. Given that you already have a fine-tuned BERT model (`finetuned\_bert\_with\_validation.pt`) and a PyTorch model (`pytorch\_model.bin`), you could combine these with new data for even more specialized training.

For the path, let's say we'll use `/content/drive/MyDrive/TrainingData/` to gather all your data. You can then preprocess and tokenize this data to feed it into your existing models for further fine-tuning.

Would you like to proceed with this idea?

## USER

Ok what kind of Files can I put in there? Answer that, but make the fucking script to do everything too.

## ASSISTANT

You can put text files like `.txt`, `.csv`, or `.json` in the `/content/drive/MyDrive/TrainingData/` directory for training. You could even use pre-tokenized data if you have it.

Here's a general outline for the script, considering you're using PyTorch and have existing models:

1. **Data Preparation**: Load and preprocess the data from `/content/drive/MyDrive/TrainingData/`.
2. **Tokenization**: Tokenize the processed text.
3. **Model Loading**: Load your existing BERT and PyTorch models.
4. **Fine-Tuning**: Fine-tune the model on the new data.
5. **Saving**: Save the newly trained model.

```
```python
import os

import torch

from transformers import BertTokenizer, BertForSequenceClassification
from torch.utils.data import DataLoader

# Data Preparation

data_path = '/content/drive/MyDrive/TrainingData/'

files = [f for f in os.listdir(data_path) if f.endswith('.txt')]

data = ""

for file in files:

    with open(os.path.join(data_path, file), 'r') as f:
```

```
data += f.read()
```

```
# Tokenization
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
```

```
tokens = tokenizer(data, padding=True, truncation=True, return_tensors='pt')
```

```
# Model Loading
```

```
bert_model_path = '/content/drive/MyDrive/000/finetuned_bert_with_validation.pt'
```

```
pytorch_model_path = '/content/drive/MyDrive/000/pytorch_model.bin'
```

```
config_path = '/content/drive/MyDrive/000/config.json'
```

```
model = BertForSequenceClassification.from_pretrained(bert_model_path, config=config_path)
```

```
model.load_state_dict(torch.load(pytorch_model_path))
```

```
# Fine-Tuning
```

```
optimizer = torch.optim.AdamW(model.parameters(), lr=1e-5)
```

```
dataloader = DataLoader(tokens, batch_size=32)
```

```
for batch in dataloader:
```

```
    optimizer.zero_grad()
```

```
    outputs = model(**batch)
```

```
    loss = outputs.loss
```

```
    loss.backward()
```

```
    optimizer.step()
```

```
# Saving
```

```
save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'

torch.save(model.state_dict(), save_path)

...
```

This script is quite simplified and assumes certain things, like your data being in text files and your using a BERT model for sequence classification. Adjustments would likely be needed based on your specific requirements.

USER

Make way better just trained on the sample data /content/sample_data There's one read me file in five CSV files

ASSISTANT

Certainly! You can fine-tune a pre-trained BERT model using sample CSV files found in `/content/sample_data/`. Below is a complete Python script that demonstrates how to do this:

1. **Data Loading**: We'll read data from the CSV files in `/content/sample_data/`.
2. **Tokenization**: We'll tokenize the text data for model training.
3. **Model Loading**: We'll use a pre-trained BERT model for fine-tuning.
4. **Fine-Tuning**: We'll fine-tune the model on the sample data.
5. **Model Saving**: Finally, we'll save the model.

```
```python
```

```
import os
```

```
import pandas as pd
```

```
import torch
```

```
from torch.utils.data import DataLoader, TensorDataset
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AdamW
```

### # 1. Data Loading

```
sample_data_path = '/content/sample_data/'
```

```
csv_files = [f for f in os.listdir(sample_data_path) if f.endswith('.csv')]
```

### # Concatenating all CSV files into one DataFrame

```
df_list = []
```

```
for csv_file in csv_files:
```

```
 df = pd.read_csv(os.path.join(sample_data_path, csv_file))
```

```
 df_list.append(df)
```

```
all_data = pd.concat(df_list, ignore_index=True)
```

### # Using a single column for demonstration (adjust as needed)

```
text_data = all_data['col1'].dropna() # Replace 'col1' with the actual column name containing text
```

### # 2. Tokenization

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
```

```
tokens = tokenizer(list(text_data), padding=True, truncation=True, return_tensors='pt')
```

```
input_ids = tokens['input_ids']
```

```
attention_mask = tokens['attention_mask']
```

### # 3. Model Loading

```
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)
```



#### # 4. Fine-Tuning

```
optimizer = AdamW(model.parameters(), lr=1e-5)

dataset = TensorDataset(input_ids, attention_mask)

dataloader = DataLoader(dataset, batch_size=32)

for batch in dataloader:

 batch_input_ids, batch_attention_mask = batch

 optimizer.zero_grad()

 outputs = model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))

 loss = outputs.loss

 loss.backward()

 optimizer.step()
```

#### # 5. Model Saving

```
save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'

torch.save(model.state_dict(), save_path)

...
```

Note: Replace ``col1`` with the actual column name that contains the text data in your CSV files. Also, the script assumes that the problem is a binary classification (`num\_labels=2`), so you may need to adjust that according to your specific case.

You'll need to run this in an environment where the required libraries are installed and where you have access to ``/content/sample\_data/`` and ``/content/drive/MyDrive/000/``.

## USER

I mean, there is fucking four CSV files in one Jason and one read me. Sorry.

## ASSISTANT

I apologize for the confusion. If you have four CSV files, one JSON file, and one README in `/content/sample_data/`, you can modify the script accordingly. Here's an updated version that caters to this specific setup:

1. **Data Loading**: Read data from the four CSV files and one JSON file in `/content/sample_data/`.
2. **Data Preparation**: Tokenize and prepare the text data.
3. **Model Loading**: Use a pre-trained BERT model for fine-tuning.
4. **Fine-Tuning**: Fine-tune the model using the sample data.
5. **Model Saving**: Save the fine-tuned model.

```
```python
```

```
import os
```

```
import pandas as pd
```

```
import json
```

```
import torch
```

```
from torch.utils.data import DataLoader, TensorDataset
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AdamW
```

```
# 1. Data Loading
```

```
sample_data_path = '/content/sample_data/'
```

```
csv_files = [f for f in os.listdir(sample_data_path) if f.endswith('.csv')]
```

```
json_file = [f for f in os.listdir(sample_data_path) if f.endswith('.json')][0]
```

```
# Concatenating all CSV files into one DataFrame
```

```
df_list = []
```

```
for csv_file in csv_files:
```

```
    df = pd.read_csv(os.path.join(sample_data_path, csv_file))
```

```
    df_list.append(df)
```

```
all_data_csv = pd.concat(df_list, ignore_index=True)
```

```
# Load JSON data
```

```
with open(os.path.join(sample_data_path, json_file), 'r') as f:
```

```
    json_data = json.load(f)
```

```
json_df = pd.DataFrame(json_data)
```

```
# Combine CSV and JSON data (Assuming both have a text column 'text_column_name')
```

```
all_data = pd.concat([all_data_csv['text_column_name'], json_df['text_column_name']],  
ignore_index=True).dropna()
```

```
# 2. Data Preparation
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
```

```
tokens = tokenizer(list(all_data), padding=True, truncation=True, return_tensors='pt')
```

```
input_ids = tokens['input_ids']
```

```
attention_mask = tokens['attention_mask']
```

```
# 3. Model Loading
```

```
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)
```

4. Fine-Tuning

```
optimizer = AdamW(model.parameters(), lr=1e-5)

dataset = TensorDataset(input_ids, attention_mask)

dataloader = DataLoader(dataset, batch_size=32)

for batch in dataloader:

    batch_input_ids, batch_attention_mask = batch

    optimizer.zero_grad()

        outputs = model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))

    loss = outputs.loss

    loss.backward()

    optimizer.step()
```

5. Model Saving

```
save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'

torch.save(model.state_dict(), save_path)

...
```

Replace ``text_column_name`` with the actual column name containing the text data in your CSV and JSON files. You may also need to adjust the number of labels (``num_labels=2``) according to your specific case.

Run this script in an environment where the required libraries are installed, and where you have access to ``/content/sample_data/`` and ``/content/drive/MyDrive/000/``.

USER

```
-----  
KeyError                                Traceback (most recent call last)  
/usr/local/lib/python3.10/dist-packages/pandas/core/indexes/base.py in get_loc(self, key, method,  
tolerance)  
    3801         try:  
-> 3802             return self._engine.get_loc(casted_key)  
    3803         except KeyError as err:
```

4 frames

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

KeyError: 'text_column_name'

The above exception was the direct cause of the following exception:

```
KeyError                                Traceback (most recent call last)  
/usr/local/lib/python3.10/dist-packages/pandas/core/indexes/base.py in get_loc(self, key, method,  
tolerance)  
    3802         return self._engine.get_loc(casted_key)  
    3803         except KeyError as err:  
-> 3804             raise KeyError(key) from err  
    3805         except TypeError:
```

```
KeyError: 'text_column_name'
```

ASSISTANT

The error suggests that the column `'text_column_name'` doesn't exist in your DataFrame. You need to replace `'text_column_name'` with the actual name of the column containing the text data in your CSV and JSON files.

Would you like to try again with the correct column name?

USER

Redo just to train on these files

ASSISTANT

To train a machine learning model on the uploaded California housing test and train CSV files, you could follow these steps:

1. **Data Loading**: Load the training and testing data from the CSV files.
2. **Data Preprocessing**: Handle missing values, normalize features, etc.
3. **Feature Engineering**: Select the features you want to use for training.
4. **Model Training**: Train a machine learning model, such as a linear regression model for this case.
5. **Model Evaluation**: Evaluate the model on the test data.
6. **Model Saving**: Save the trained model.

Let's start by examining the first few rows of both the training and test datasets.

ASSISTANT

It looks like there was an issue with the code execution. Let me try loading the data again.

ASSISTANT

I apologize for the inconvenience. It seems there's an issue with the code environment. Let me try to resolve it and load the data again.

ASSISTANT

I apologize for the ongoing issues. It appears that the variable for the data path got lost. Let's try to reload the data one more time.

ASSISTANT

Great, the data has been successfully loaded. The training and testing datasets have the following columns:

- `longitude`
- `latitude`
- `housing_median_age`
- `total_rooms`
- `total_bedrooms`
- `population`
- `households`
- `median_income`
- `median_house_value`

Would you like to proceed with a specific type of machine learning model, or should I go ahead and choose one for you?

USER

```
import os
```

```
import pandas as pd
```

```
import json
```

```
import torch
```

```
from torch.utils.data import DataLoader, TensorDataset
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AdamW
```

1. Data Loading

```
sample_data_path = '/content/sample_data/'
```

```
csv_files = [f for f in os.listdir(sample_data_path) if f.endswith('.csv')]
```

```
json_file = [f for f in os.listdir(sample_data_path) if f.endswith('.json')][0]
```

Concatenating all CSV files into one DataFrame

```
df_list = []
```

```
for csv_file in csv_files:
```

```
    df = pd.read_csv(os.path.join(sample_data_path, csv_file))
```

```
    df_list.append(df)
```

```
all_data_csv = pd.concat(df_list, ignore_index=True)
```

Load JSON data

```
with open(os.path.join(sample_data_path, json_file), 'r') as f:
```



```

json_data = json.load(f)

json_df = pd.DataFrame(json_data)

# Combine CSV and JSON data (Assuming both have a text column 'text_column_name')
all_data = pd.concat([all_data_csv['text_column_name'], json_df['text_column_name']],
ignore_index=True).dropna()

# 2. Data Preparation

tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

tokens = tokenizer(list(all_data), padding=True, truncation=True, return_tensors='pt')

input_ids = tokens['input_ids']

attention_mask = tokens['attention_mask']

# 3. Model Loading

model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)

# 4. Fine-Tuning

optimizer = AdamW(model.parameters(), lr=1e-5)

dataset = TensorDataset(input_ids, attention_mask)

dataloader = DataLoader(dataset, batch_size=32)

for batch in dataloader:

    batch_input_ids, batch_attention_mask = batch

    optimizer.zero_grad()

    outputs = model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))

```

```
loss = outputs.loss
```

```
loss.backward()
```

```
optimizer.step()
```

5. Model Saving

```
save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'
```

```
torch.save(model.state_dict(), save_path)
```

ASSISTANT

The script you've provided is a Python script that combines several aspects of your end-to-end process, from data loading to model training and saving. It looks like it's designed to:

1. Load data from multiple CSV files and a JSON file from a specified directory (`/content/sample_data/`).
2. Concatenate these files into a single DataFrame.
3. Tokenize the text data using a BERT tokenizer.
4. Load a pre-trained BERT model for sequence classification.
5. Fine-tune this model on the loaded text data.
6. Save the fine-tuned model to a specified path (`/content/drive/MyDrive/000/new_finetuned_model.pt`).

However, there are a few things to note:

- The script assumes that both the CSV and JSON files have a text column named `'text_column_name'`, which you'll need to modify according to your actual column names.

- You're using BERT for sequence classification with only two labels. Make sure your task aligns with this.
- The script saves the fine-tuned model to a path; make sure that path exists or is writable.

Would you like me to modify the script based on specific requirements, or do you have any questions about it?

USER

Combine

ASSISTANT

To combine your provided Python script with the end-to-end process you're interested in, I've made a few changes to align it with your needs:

1. I've updated the script to work with the sample data files you've uploaded: ``/mnt/data/california_housing_test.csv`` and ``/mnt/data/california_housing_train.csv``. These are used as placeholders for your real data.
2. The script assumes a two-class classification problem, so please adjust ``num_labels`` according to your specific use case.
3. The fine-tuned model is saved to ``/content/drive/MyDrive/000/new_finetuned_model.pt``.

Here's the combined script:

```
```python
import os

import pandas as pd
```

```
import torch
```

```
from torch.utils.data import DataLoader, TensorDataset
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AdamW
```

### # 1. Data Loading

```
Update the paths to the files you've uploaded
```

```
train_data_path = '/mnt/data/california_housing_train.csv'
```

```
test_data_path = '/mnt/data/california_housing_test.csv'
```

```
Load CSV files into DataFrames
```

```
train_df = pd.read_csv(train_data_path)
```

```
test_df = pd.read_csv(test_data_path)
```

```
Combine train and test data (Assuming both have a text column 'text_column_name')
```

```
Replace 'text_column_name' with the actual name of the text column you're interested in
```

```
all_data = pd.concat([train_df['text_column_name'], test_df['text_column_name']],
ignore_index=True).dropna()
```

### # 2. Data Preparation

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
```

```
tokens = tokenizer(list(all_data), padding=True, truncation=True, return_tensors='pt')
```

```
input_ids = tokens['input_ids']
```

```
attention_mask = tokens['attention_mask']
```

### # 3. Model Loading

```
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)
```

#### # 4. Fine-Tuning

```
optimizer = AdamW(model.parameters(), lr=1e-5)

dataset = TensorDataset(input_ids, attention_mask)

dataloader = DataLoader(dataset, batch_size=32)

for batch in dataloader:

 batch_input_ids, batch_attention_mask = batch

 optimizer.zero_grad()

 outputs = model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))

 loss = outputs.loss

 loss.backward()

 optimizer.step()
```

#### # 5. Model Saving

```
Update this path to where you'd like to save the model

save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'

torch.save(model.state_dict(), save_path)

...
```

Replace ``text\_column\_name`` with the actual column name that contains the text data you wish to use for training. If your data is not primarily text and you're looking to do something other than text classification, the script would need to be adapted accordingly.

## USER

No, I want to train my Models on it Gave you was fucked up because it didn't have the column names Wanted you to complete the whole fucking thing with onlimport os

```
import pandas as pd
```

```
import json
```

```
import torch
```

```
from torch.utils.data import DataLoader, TensorDataset
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AdamW
```

```
1. Data Loading
```

```
sample_data_path = '/content/sample_data/'
```

```
csv_files = [f for f in os.listdir(sample_data_path) if f.endswith('.csv')]
```

```
json_file = [f for f in os.listdir(sample_data_path) if f.endswith('.json')][0]
```

```
Concatenating all CSV files into one DataFrame
```

```
df_list = []
```

```
for csv_file in csv_files:
```

```
 df = pd.read_csv(os.path.join(sample_data_path, csv_file))
```

```
 df_list.append(df)
```

```
all_data_csv = pd.concat(df_list, ignore_index=True)
```

```
Load JSON data
```

```
with open(os.path.join(sample_data_path, json_file), 'r') as f:
```

```
 json_data = json.load(f)
```

```
json_df = pd.DataFrame(json_data)
```

```
Combine CSV and JSON data (Assuming both have a text column 'text_column_name')
```

```
all_data = pd.concat([all_data_csv['text_column_name'], json_df['text_column_name']],
ignore_index=True).dropna()
```

## # 2. Data Preparation

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
tokens = tokenizer(list(all_data), padding=True, truncation=True, return_tensors='pt')
input_ids = tokens['input_ids']
attention_mask = tokens['attention_mask']
```

## # 3. Model Loading

```
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)
```

## # 4. Fine-Tuning

```
optimizer = AdamW(model.parameters(), lr=1e-5)
dataset = TensorDataset(input_ids, attention_mask)
dataloader = DataLoader(dataset, batch_size=32)
```

for batch in dataloader:

```
 batch_input_ids, batch_attention_mask = batch

 optimizer.zero_grad()

 outputs = model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))

 loss = outputs.loss

 loss.backward()

 optimizer.step()
```

## # 5. Model Saving

```
save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'
```

```
torch.save(model.state_dict(), save_path)
```

## ASSISTANT

I apologize for the confusion. Let's correct this. The script you've provided is almost complete but assumes there's a text column ('text\_column\_name') in your data. Since the data files you uploaded ('california\_housing\_test.csv' and 'california\_housing\_train.csv') seem to be about California housing, they're unlikely to contain text data for text-based machine learning.

If you want to train a BERT model on non-text data, that's a different kind of problem that involves feature extraction and other preprocessing steps.

Could you please clarify what exactly you want to do with these specific datasets? Are you trying to predict a specific column using text data? Or do you have text data in another dataset that you would like to use for training the BERT model?

Once I have more details, I can help you create a complete script tailored to your needs.

## USER

Whatever I can do to use both of my models that I gave you and train another one with all of them together

## ASSISTANT

If you'd like to combine the two pre-trained models ('finetuned\_bert\_with\_validation.pt' and



`pytorch\_model.bin` with its associated `config.json`) and train a new one using the California housing data, it gets a bit complex. Combining two different pre-trained models into a single new model isn't straightforward because each model has its own architecture and learned parameters.

Here's a general approach you could take:

1. **Load both models**: Use PyTorch to load both models into memory.
2. **Architecture Decision**: Decide what the architecture of your new model will look like. Will it be an ensemble of both models? Will you take layers from one and add them to another? This part is crucial and may require some expertise in machine learning model architectures.
3. **Data Preparation**: Prepare your California housing data for training. Since it's not text data, you'll need to decide how to feed this into your model(s).
4. **Training Loop**: Train your new model architecture on your training data.
5. **Evaluation**: Evaluate the new model on some validation data.
6. **Save the Model**: Save the newly trained model to disk.

Would you like to proceed with this approach, or do you have another idea in mind?

## USER

I want a joblib model

## ASSISTANT

If you're looking to create a new model using scikit-learn and save it as a joblib file, you can certainly do that. Since you've uploaded the California housing test and train datasets, we can use them to train a simple machine learning model and then save it using joblib.

Here's how to do it:

1. Load the California housing data.
2. Preprocess the data.
3. Train a scikit-learn model.
4. Save the model using joblib.

Would you like to proceed with this?

**## USER**

But I want to convert my other models together to convert it to the jobLib

**## ASSISTANT**

Combining models trained with different libraries (like PyTorch for your BERT models and scikit-learn for a joblib model) into a single joblib file can be complex. Joblib is generally for scikit-learn models and won't directly handle PyTorch models.

However, one approach could be to wrap your PyTorch models and any scikit-learn models into a custom Python class. This class could handle the pre-processing and predictions for each model internally. Once that's done, you could then save this custom class using joblib.

Would you like to see an example code snippet for this approach?

**## USER**

Do whatever you think is best to make it work I don't wanna have no fucking err

**## ASSISTANT**

Alright, let's try a step-by-step approach to get this right without errors.

1. **\*\*Load Pre-trained Models\*\***: Load both of your existing PyTorch models.
2. **\*\*Data Preprocessing\*\***: Preprocess the new training data in the format required for these models.
3. **\*\*Model Training\*\***: Train the new model on this preprocessed data.
4. **\*\*Ensemble\*\***: Create a custom Python class to combine these three models. This class will take an input, make predictions using each of the three models, and return an ensemble prediction.
5. **\*\*Save as Joblib\*\***: Save this ensemble model as a joblib file.

I can't run PyTorch code here, but I can provide you with code snippets to achieve these steps.

Would that work for you?

**## USER**

```
import os
```

```
import pandas as pd
```

```
import json
```

```
import torch
```

```
from torch.utils.data import DataLoader, TensorDataset
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AdamW
```

```
1. Data Loading
```

```
sample_data_path = '/content/sample_data/'
```

```
csv_files = [f for f in os.listdir(sample_data_path) if f.endswith('.csv')]
```

```
json_file = [f for f in os.listdir(sample_data_path) if f.endswith('.json')][0]
```

```
Concatenating all CSV files into one DataFrame
```

```

df_list = []

for csv_file in csv_files:

 df = pd.read_csv(os.path.join(sample_data_path, csv_file))

 df_list.append(df)

all_data_csv = pd.concat(df_list, ignore_index=True)

Load JSON data

with open(os.path.join(sample_data_path, json_file), 'r') as f:

 json_data = json.load(f)

json_df = pd.DataFrame(json_data)

Combine CSV and JSON data (Assuming both have a text column 'text_column_name')

all_data = pd.concat([all_data_csv['text_column_name'], json_df['text_column_name']],
ignore_index=True).dropna()

2. Data Preparation

tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

tokens = tokenizer(list(all_data), padding=True, truncation=True, return_tensors='pt')

input_ids = tokens['input_ids']

attention_mask = tokens['attention_mask']

3. Model Loading

model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)

4. Fine-Tuning

optimizer = AdamW(model.parameters(), lr=1e-5)

```

```

dataset = TensorDataset(input_ids, attention_mask)

dataloader = DataLoader(dataset, batch_size=32)

for batch in dataloader:

 batch_input_ids, batch_attention_mask = batch

 optimizer.zero_grad()

 outputs = model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))

 loss = outputs.loss

 loss.backward()

 optimizer.step()

```

## # 5. Model Saving

```

save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'

torch.save(model.state_dict(), save_path)

```

Make that do whatever we just plan on. Convert that to the new one where we only do the two CSV files.

## ## ASSISTANT

Certainly, you can modify the code to just use the two CSV files that you uploaded for training. Given that you've uploaded `california\_housing\_test.csv` and `california\_housing\_train.csv`, you can adjust the code as follows:

1. Update the data loading section to only use these specific CSV files.
2. Replace `text\_column\_name` with the actual column name in these CSV files that you want to use for training.

Here's how the modified code would look:

```
```python

import os

import pandas as pd

import torch

from torch.utils.data import DataLoader, TensorDataset

from transformers import BertTokenizer, BertForSequenceClassification, AdamW

# 1. Data Loading

data_path = '/mnt/data/' # Update this to the path where the uploaded files are

csv_files = ['california_housing_test.csv', 'california_housing_train.csv']

# Concatenating both CSV files into one DataFrame

df_list = []

for csv_file in csv_files:

    df = pd.read_csv(os.path.join(data_path, csv_file))

    df_list.append(df)

all_data_csv = pd.concat(df_list, ignore_index=True)

# Update the column name to the column you want to use for training

target_column = 'your_target_column_here'

all_data = all_data_csv[target_column].dropna()

# 2. Data Preparation
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

tokens = tokenizer(list(all_data), padding=True, truncation=True, return_tensors='pt')

input_ids = tokens['input_ids']

attention_mask = tokens['attention_mask']
```

3. Model Loading

```
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)
```

4. Fine-Tuning

```
optimizer = AdamW(model.parameters(), lr=1e-5)

dataset = TensorDataset(input_ids, attention_mask)

dataloader = DataLoader(dataset, batch_size=32)
```

for batch in dataloader:

```
    batch_input_ids, batch_attention_mask = batch

    optimizer.zero_grad()

        outputs = model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))

    loss = outputs.loss

    loss.backward()

    optimizer.step()
```

5. Model Saving

```
save_path = '/content/drive/MyDrive/000/new_finetuned_model.pt'

torch.save(model.state_dict(), save_path)
```

```
...
```

Replace ``your_target_column_here`` with the column name you want to use for training from the CSV files. This script will load the data, preprocess it, train a new model, and save it.

USER

Load Pre-trained Models: Load both of your existing PyTorch models.

Data Preprocessing: Preprocess the new training data in the format required for these models.

Model Training: Train the new model on this preprocessed data.

Ensemble: Create a custom Python class to combine these three models. This class will take an input, make predictions using each of the three models, and return an ensemble prediction.

Save as Joblib: Save this ensemble model as a joblib file. But I'm in collaboratory, so use the fucking content directory. That's where the CSV files are stored at and I already gave you the paths where my models are located and I configuration is.

ASSISTANT

Alright, let's break this down step-by-step:

1. **Load Pre-trained Models**: Load both of your existing PyTorch models.
2. **Data Preprocessing**: Preprocess the data in the format required for these models.
3. **Model Training**: Train a new model on this preprocessed data.
4. **Ensemble**: Combine these three models into one ensemble model.
5. **Save as Joblib**: Save this ensemble model as a joblib file.

Here's how you can implement all these steps in code:

```
```python
```



```
import os

import pandas as pd

import torch

from torch.utils.data import DataLoader, TensorDataset

from transformers import BertTokenizer, BertForSequenceClassification, AdamW

import joblib
```

## # 1. Load Pre-trained Models

```
model1 = BertForSequenceClassification.from_pretrained('/content/drive/MyDrive/000/finetuned_bert_with_validation.pt')

model2 = BertForSequenceClassification.from_pretrained('/content/drive/MyDrive/000/pytorch_model.bin', config='/content/drive/MyDrive/000/config.json')
```

## # 2. Data Preprocessing

```
data_path = '/content/sample_data/' # Update this to the path where the sample_data is

csv_files = ['california_housing_test.csv', 'california_housing_train.csv']
```

### # Load and concatenate CSV files

```
df_list = []

for csv_file in csv_files:

 df = pd.read_csv(os.path.join(data_path, csv_file))

 df_list.append(df)

all_data_csv = pd.concat(df_list, ignore_index=True)
```

```
Update the column name to the column you want to use for training
```

```
target_column = 'your_target_column_here'
```

```
all_data = all_data_csv[target_column].dropna()
```

```
Tokenize the data
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
```

```
tokens = tokenizer(list(all_data), padding=True, truncation=True, return_tensors='pt')
```

```
input_ids = tokens['input_ids']
```

```
attention_mask = tokens['attention_mask']
```

```
3. Model Training - New Model
```

```
new_model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)
```

```
optimizer = AdamW(new_model.parameters(), lr=1e-5)
```

```
dataset = TensorDataset(input_ids, attention_mask)
```

```
dataloader = DataLoader(dataset, batch_size=32)
```

```
Training loop
```

```
for batch in dataloader:
```

```
 batch_input_ids, batch_attention_mask = batch
```

```
 optimizer.zero_grad()
```

```
 outputs = new_model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
labels=torch.zeros(batch_input_ids.shape[0], dtype=torch.long))
```

```
 loss = outputs.loss
```

```
 loss.backward()
```

```
 optimizer.step()
```

#### # 4. Ensemble

```
class EnsembleModel:
```

```
 def __init__(self, model1, model2, new_model):
```

```
 self.model1 = model1
```

```
 self.model2 = model2
```

```
 self.new_model = new_model
```

```
 def predict(self, input_ids, attention_mask):
```

```
 with torch.no_grad():
```

```
 pred1 = self.model1(input_ids=input_ids, attention_mask=attention_mask).logits
```

```
 pred2 = self.model2(input_ids=input_ids, attention_mask=attention_mask).logits
```

```
 pred_new = self.new_model(input_ids=input_ids, attention_mask=attention_mask).logits
```

```
 ensemble_pred = (pred1 + pred2 + pred_new) / 3
```

```
 return torch.argmax(ensemble_pred, dim=1)
```

```
Create Ensemble Model
```

```
ensemble_model = EnsembleModel(model1, model2, new_model)
```

#### # 5. Save as Joblib

```
joblib_file = "/content/ensemble_model.joblib"
```

```
joblib.dump(ensemble_model, joblib_file)
```

```
...
```

Replace ``your\_target\_column\_here`` with the actual column name from the CSV files you want to use for training. This script should do what you're looking for.

## USER

Replace your target column here with the actual column name of the CSV files that I just fucking gave you and make everything

## ASSISTANT

I understand you're looking to create a pipeline that utilizes your existing models and new data to train an ensemble model. Let's break this down step-by-step.

1. **\*\*Load Pre-trained Models\*\***: You've provided paths to two pre-trained models (`/content/drive/MyDrive/000/finetuned_bert_with_validation.pt`` and `/content/drive/MyDrive/000/pytorch_model.bin``) along with their configuration (`/content/drive/MyDrive/000/config.json``). These will be loaded first.
2. **\*\*Data Preprocessing\*\***: The new data you've uploaded (`california_housing_test.csv`` and `california_housing_train.csv``) will be preprocessed. The target column for these files is generally "median\_house\_value."
3. **\*\*Model Training\*\***: A new model will be trained on this preprocessed data.
4. **\*\*Ensemble\*\***: An ensemble model will be created that combines the predictions from the two pre-trained models and the newly trained model.
5. **\*\*Save as Joblib\*\***: Finally, this ensemble model will be saved as a joblib file in the `/content`` directory.

Given the complexity of this task, it will take some time. Let's start with loading your pre-trained

models. Would you like to proceed?

## ## USER

Go

## ## ASSISTANT

It appears that the `transformers` library is not available in this environment. Since I can't install packages, you'd need to install it yourself in your local Colab environment using `!pip install transformers` to proceed. Would you like to continue with the next steps after installing the package?

## ## USER

RuntimeError      Traceback (most recent call last)

```
<ipython-input-7-088ce2578315> in <cell line: 17>()
```

15

16 # Load the pre-trained model weights

```
---> 17 finetuned_bert_model.load_state_dict(torch.load(finetuned_bert_path))
```

```
18 pytorch_model.load_state_dict(torch.load(pytorch_model_path))
```

19

```

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in load_state_dict(self,
state_dict, strict)

```

2039

```
2040 if len(error_msgs) > 0:
```

```
-> 2041 raise RuntimeError('Error(s) in loading state_dict for {}'.format(
```

```
2042 self.__class__.__name__, "\n\t".join(error_msgs)))
```

```
2043 return _IncompatibleKeys(missing_keys, unexpected_keys)
```

RuntimeError: Error(s) in loading state\_dict for BertModel:

Missing	key(s)	in	state_dict:	
			"embeddings.word_embeddings.weight",	
	"embeddings.position_embeddings.weight",		"embeddings.token_type_embeddings.weight",	
	"embeddings.LayerNorm.weight",		"embeddings.LayerNorm.bias",	
	"encoder.layer.0.attention.self.query.weight",		"encoder.layer.0.attention.self.query.bias",	
	"encoder.layer.0.attention.self.key.weight",		"encoder.layer.0.attention.self.key.bias",	
	"encoder.layer.0.attention.self.value.weight",		"encoder.layer.0.attention.self.value.bias",	
	"encoder.layer.0.attention.output.dense.weight",		"encoder.layer.0.attention.output.dense.bias",	
	"encoder.layer.0.attention.output.LayerNorm.weight",			
	"encoder.layer.0.attention.output.LayerNorm.bias",		"encoder.layer.0.intermediate.dense.weight",	
	"encoder.layer.0.intermediate.dense.bias",		"encoder.layer.0.output.dense.weight",	
	"encoder.layer.0.output.dense.bias",		"encoder.layer.0.output.LayerNorm.weight",	
	"encoder.layer.0.output.LayerNorm.bias",		"encoder.layer.1.attention.self.query.weight",	
	"encoder.layer.1.attention.self.query.bias",		"encoder.layer.1.attention.self.key.weight",	
	"encoder.layer.1.attention.self.key.bias",		"encoder.layer.1.attention.self.value.weight",	
	"encoder.layer.1.attention.self.value.bias",		"encoder.layer.1.attention.output.dense.weight",	
	"encoder.layer.1.attention.output.dense.bias",		"encoder.layer.1.attention.output.LayerNorm.weight",	
	"encoder.layer.1.attention.output.LayerNorm.bias",		"encoder.layer.1.intermediate.dense.weight",	
	"encoder.layer.1.intermediate.dense.bias",		"enco...	

Unexpected	key(s)	in	state_dict:	
			"bert.embeddings.word_embeddings.weight",	
	"bert.embeddings.position_embeddings.weight",		"bert.embeddings.token_type_embeddings.weight",	
	"bert.embeddings.LayerNorm.weight",		"bert.embeddings.LayerNorm.bias",	
	"bert.encoder.layer.0.attention.self.query.weight",		"bert.encoder.layer.0.attention.self.query.bias",	

"bert.encoder.layer.0.attention.self.key.weight", "bert.encoder.layer.0.attention.self.key.bias",  
"bert.encoder.layer.0.attention.self.value.weight", "bert.encoder.layer.0.attention.self.value.bias",  
"bert.encoder.layer.0.attention.output.dense.weight",  
"bert.encoder.layer.0.attention.output.dense.bias",  
"bert.encoder.layer.0.attention.output.LayerNorm.weight",  
"bert.encoder.layer.0.attention.output.LayerNorm.bias",  
"bert.encoder.layer.0.intermediate.dense.weight", "bert.encoder.layer.0.intermediate.dense.bias",  
"bert.encoder.layer.0.output.dense.weight", "bert.encoder.layer.0.output.dense.bias",  
"bert.encoder.layer.0.output.LayerNorm.weight", "bert.encoder.layer.0.output.LayerNorm.bias",  
"bert.encoder.layer.1.attention.self.query.weight", "bert.encoder.layer.1.attention.self.query.bias",  
"bert.encoder.layer.1.attention.self.key.weight", "bert.encoder.layer.1.attention.self.key.bias",  
"bert.encoder.layer.1.attention.self.value.weight", "bert.encoder.layer.1.attention.self.value.bias",  
"bert.encoder.layer.1.attention.output.dense.weight",  
"bert.encoder.layer.1.attention.output.dense.bias", "bert.encoder.layer.1.attention.output.LayerNorm

## ASSISTANT

The error indicates that the keys in the saved state\_dict do not match the keys in the model you're trying to load. This could happen if the architecture of the model you're loading differs from the architecture of the saved model. Make sure both are the same.

Also, if you're loading a model trained on a specific task (like Sequence Classification) into a general BERT model, you'd need to adjust the keys to remove task-specific prefixes (like `bert.`).

Would you like to proceed with adjusting the keys to make them compatible?

## USER

Y

## ASSISTANT

To adjust the keys and make them compatible, we can strip the task-specific prefixes (like `bert.`) from the state dictionary of the saved model. Then we can try loading it again.

Let's start by adjusting the keys for the saved finetuned BERT model and the PyTorch model. After that, we can proceed with the ensemble training using the California housing CSV files you've uploaded.

Here's how you can adjust the keys:

```
```python
```

```
# Function to fix the keys
```

```
def fix_state_dict(state_dict):
```

```
    new_state_dict = {}
```

```
    for k, v in state_dict.items():
```

```
        name = k
```

```
        if name.startswith('bert.')
```

```
            name = name[5:] # remove 'bert.' prefix
```

```
            new_state_dict[name] = v
```

```
    return new_state_dict
```

```
# Load the original state_dict
```

```
finetuned_bert_state_dict
```

=

```
torch.load('/content/drive/MyDrive/000/finetuned_bert_with_validation.pt')
```



```

pytorch_model_state_dict = torch.load('/content/drive/MyDrive/000/pytorch_model.bin')

# Fix the keys in state_dict

fixed_finetuned_bert_state_dict = fix_state_dict(finetuned_bert_state_dict)

fixed_pytorch_model_state_dict = fix_state_dict(pytorch_model_state_dict)


# Load the model again

finetuned_bert_model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
num_labels=2)

finetuned_bert_model.load_state_dict(fixed_finetuned_bert_state_dict)


pytorch_model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
num_labels=2)

pytorch_model.load_state_dict(fixed_pytorch_model_state_dict)
...

```

Once the keys are adjusted and the models are loaded, you can proceed with training the ensemble model. Would you like to continue with that?

USER

Downloading (?)lve/main/config.json: 100%

570/570 [00:00<00:00, 39.5kB/s]

Downloading model.safetensors: 100%

440M/440M [00:01<00:00, 364MB/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
-----

RuntimeError                                Traceback (most recent call last)
<ipython-input-8-b8047e31d9d0> in <cell line: 21>()
    19 # Load the model again
    20 finetuned_bert_model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
num_labels=2)
--> 21 finetuned_bert_model.load_state_dict(fixed_finetuned_bert_state_dict)
    22
    23 pytorch_model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
num_labels=2)

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py      in      load_state_dict(self,
state_dict, strict)
    2039
    2040     if len(error_msgs) > 0:
-> 2041         raise RuntimeError('Error(s) in loading state_dict for {}:'\n\t{}'.format(
    2042             self.__class__.__name__, "\n\t".join(error_msgs)))
    2043     return _IncompatibleKeys(missing_keys, unexpected_keys)

RuntimeError: Error(s) in loading state_dict for BertForSequenceClassification:
Missing      key(s)      in      state_dict:      "bert.embeddings.word_embeddings.weight",
```

"bert.embeddings.position_embeddings.weight", "bert.embeddings.token_type_embeddings.weight",
"bert.embeddings.LayerNorm.weight", "bert.embeddings.LayerNorm.bias",
"bert.encoder.layer.0.attention.self.query.weight", "bert.encoder.layer.0.attention.self.query.bias",
"bert.encoder.layer.0.attention.self.key.weight", "bert.encoder.layer.0.attention.self.key.bias",
"bert.encoder.layer.0.attention.self.value.weight", "bert.encoder.layer.0.attention.self.value.bias",
"bert.encoder.layer.0.attention.output.dense.weight",
"bert.encoder.layer.0.attention.output.dense.bias",
"bert.encoder.layer.0.attention.output.LayerNorm.weight",
"bert.encoder.layer.0.attention.output.LayerNorm.bias",
"bert.encoder.layer.0.intermediate.dense.weight", "bert.encoder.layer.0.intermediate.dense.bias",
"bert.encoder.layer.0.output.dense.weight", "bert.encoder.layer.0.output.dense.bias",
"bert.encoder.layer.0.output.LayerNorm.weight", "bert.encoder.layer.0.output.LayerNorm.bias",
"bert.encoder.layer.1.attention.self.query.weight", "bert.encoder.layer.1.attention.self.query.bias",
"bert.encoder.layer.1.attention.self.key.weight", "bert.encoder.layer.1.attention.self.key.bias",
"bert.encoder.layer.1.attention.self.value.weight", "bert.encoder.layer.1.attention.self.value.bias",
"bert.encoder.layer.1.attention.output.dense.weight",
"bert.encoder.layer.1.attention.output.dense.bias",
"bert.encoder.layer.1.attention.output.LayerNorm.wei...

Unexpected key(s) in state_dict: "embeddings.word_embeddings.weight",
"embeddings.position_embeddings.weight", "embeddings.token_type_embeddings.weight",
"embeddings.LayerNorm.weight", "embeddings.LayerNorm.bias",
"encoder.layer.0.attention.self.query.weight", "encoder.layer.0.attention.self.query.bias",
"encoder.layer.0.attention.self.key.weight", "encoder.layer.0.attention.self.key.bias",
"encoder.layer.0.attention.self.value.weight", "encoder.layer.0.attention.self.value.bias",
"encoder.layer.0.attention.output.dense.weight", "encoder.layer.0.attention.output.dense.bias",
"encoder.layer.0.attention.output.LayerNorm.weight",

"encoder.layer.0.attention.output.LayerNorm.bias",	"encoder.layer.0.intermediate.dense.weight",
"encoder.layer.0.intermediate.dense.bias",	"encoder.layer.0.output.dense.weight",
"encoder.layer.0.output.dense.bias",	"encoder.layer.0.output.LayerNorm.weight",
"encoder.layer.0.output.LayerNorm.bias",	"encoder.layer.1.attention.self.query.weight",
"encoder.layer.1.attention.self.query.bias",	"encoder.layer.1.attention.self.key.weight",
"encoder.layer.1.attention.self.key.bias",	"encoder.layer.1.attention.self.value.weight",
"encoder.layer.1.attention.self.value.bias",	"encoder.layer.1.attention.output.dense.weight",
"encoder.layer.1.attention.output.dense.bias",	"encoder.layer.1.attention.output.LayerNorm.weight",
"encoder.layer.1.attention.output.LayerNorm.bias",	"encoder.layer.1.intermediate.dense.weight",
"encoder.layer.1.intermediate.dense.bias",	"e...

ASSISTANT

It looks like there's still a mismatch between the keys in the state dictionaries of your saved models and the newly initialized BertForSequenceClassification models. This mismatch is likely because the saved models have different architecture layers or are of a different type altogether.

To successfully load your saved models, we need to make sure that:

1. The model architecture in the code matches the architecture of the saved models.
2. The keys in the state dictionary match.

Would you be able to confirm what specific model architectures were used for the saved models?

This will help me in providing a more accurate solution.

USER

PyTorch

ASSISTANT

Great, thanks for confirming that the models are in PyTorch format. Now, the next step is to adjust the code to load your specific models correctly and proceed with the ensemble training.

Here's a rough plan:

1. Load your existing PyTorch models (`finetuned_bert_with_validation.pt` and `pytorch_model.bin`) and their configurations (`config.json`).
2. Preprocess the new training data from the uploaded CSV files.
3. Train a new model on this preprocessed data.
4. Create an ensemble model that uses all three models to make predictions.
5. Save this ensemble model as a Joblib file in the `content` directory.

Let's start by loading your existing models. Could you confirm whether both of your existing models are BERT models or if they are of different architectures? This will help me in loading them correctly.

USER

Give me the script

ASSISTANT

Certainly. Below is the Python script that accomplishes the following tasks:

1. Loads two pre-existing PyTorch models (`finetuned_bert_with_validation.pt` and `pytorch_model.bin`) and their configuration (`config.json`).
2. Preprocesses the new training data from the uploaded CSV files (`california_housing_test.csv`)

and `california_housing_train.csv`).

3. Trains a new model on this preprocessed data.
4. Creates an ensemble model that utilizes all three models to make predictions.
5. Saves this ensemble model as a Joblib file.

The code assumes that your existing models and the new CSV data are in the directories specified. The target column in the CSV files is assumed to be 'median_house_value', and the feature columns are other numerical columns.

Note: This script is meant for conceptual guidance, and you may need to adjust the paths and parameters to match your specific setup.

```
```python
import os

import pandas as pd

from sklearn.ensemble import VotingClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

import joblib

import torch

from transformers import BertForSequenceClassification

Load existing PyTorch models

finetuned_bert_path = '/content/drive/MyDrive/000/finetuned_bert_with_validation.pt'
```

```
pytorch_model_path = '/content/drive/MyDrive/000/pytorch_model.bin'

finetuned_bert_model = BertForSequenceClassification.from_pretrained('bert-base-uncased')

pytorch_model = BertForSequenceClassification.from_pretrained('bert-base-uncased')

finetuned_bert_model.load_state_dict(torch.load(finetuned_bert_path))

pytorch_model.load_state_dict(torch.load(pytorch_model_path))
```

```
Load new training data
```

```
csv_file1 = '/content/california_housing_train.csv'

csv_file2 = '/content/california_housing_test.csv'

df1 = pd.read_csv(csv_file1)

df2 = pd.read_csv(csv_file2)

data = pd.concat([df1, df2])
```

```
Preprocessing
```

```
X = data.drop(['median_house_value'], axis=1)

y = data['median_house_value']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)
```

```
Train a new Sklearn model
```

```
lr_model = LinearRegression()

lr_model.fit(X_train_scaled, y_train)
```

```
Ensemble model (conceptual, actual implementation will vary)
```

```

class EnsembleModel:

 def __init__(self, model1, model2, model3):

 self.model1 = model1

 self.model2 = model2

 self.model3 = model3

 def predict(self, X):

 pred1 = self.model1.predict(X)

 pred2 = self.model2.predict(X)

 pred3 = self.model3.predict(X)

 final_pred = (pred1 + pred2 + pred3) / 3

 return final_pred

ensemble_model = EnsembleModel(finetuned_bert_model, pytorch_model, lr_model)

Save ensemble model

joblib_file = "/content/ensemble_model.pkl"

joblib.dump(ensemble_model, joblib_file)

Test ensemble model

ensemble_predictions = ensemble_model.predict(X_test_scaled)

mse = mean_squared_error(y_test, ensemble_predictions)

print(f'Mean Squared Error of Ensemble Model: {mse}')

...

```

This is a simplified and conceptual example. The actual implementation of the ensemble model



would depend on how you want to combine the predictions from the three models. Also, the code for loading PyTorch models and making predictions is simplified for demonstration purposes.

## USER

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

-----

AttributeError Traceback (most recent call last)

<ipython-input-10-ba91885381bf> in <cell line: 60>()

58

59 # Test ensemble model

---> 60 ensemble\_predictions = ensemble\_model.predict(X\_test\_scaled)

61 mse = mean\_squared\_error(y\_test, ensemble\_predictions)

62 print(f'Mean Squared Error of Ensemble Model: {mse}')

1 frames

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in \_\_getattr\_\_(self, name)

1612 if name in modules:

1613 return modules[name]

-> 1614 raise AttributeError("{}' object has no attribute '{}'".format(

```
1615 type(self).__name__, name))
```

```
1616
```

AttributeError: 'BertForSequenceClassification' object has no attribute 'predict'

## ASSISTANT

The error arises because the `BertForSequenceClassification` PyTorch models don't have a `predict` method by default. In the ensemble model code, I used the `predict` method as if it were an sklearn model. PyTorch models usually use a `forward` method for making predictions.

Moreover, the Bert models and the linear regression model aren't compatible out-of-the-box for ensembling like this because they expect different kinds of input and output different kinds of results (one is for NLP, the other is for regression on numerical data).

Given your requirements, a more complex approach would be required to effectively combine these models into a single ensemble model. This would involve manually writing the prediction logic for each model in the ensemble and determining how to weight or combine these predictions into a final prediction.

However, the real question here is whether it makes sense to ensemble these models in the first place. Bert models are designed for text classification or similar NLP tasks, while the linear regression model is designed for numerical regression tasks. If your new data involves both text and numerical features, a more involved feature engineering and modeling approach would be necessary.

Would you like to proceed with a more complex example or clarify the kind of data you're working

with?

## USER

Just do whatever you think is the best

## ASSISTANT

To create an ensemble of models, especially when they are fundamentally different like Bert models and linear regression models, you need to ensure that the data being passed to each is appropriate and that their outputs can be meaningfully combined.

However, given the data and models you mentioned, it's a bit unclear how to combine a language model like Bert with a numerical model trained on housing data. The tasks they perform are quite different.

So, I'll demonstrate the following steps:

1. Preprocess the California housing data (as it's numerical) for a linear regression model.
2. Preprocess text data (I'll generate some dummy text data since no specific text data was provided) for Bert models.
3. Train a simple linear regression model on the California housing data.
4. Load your pre-trained Bert models.
5. Create an ensemble model to combine these.

Since this is a simplified example, the ensemble model will just average the predictions of the three models (which doesn't make a lot of practical sense given that they are for different tasks, but should illustrate the process).

Let's start.

## ASSISTANT

It looks like there was an issue with importing the PyTorch library, possibly due to the environment limitations. Normally, you should be able to run this code in a local environment or on platforms like Google Colab without this issue. Would you like me to provide you with the code so you can run it on your end?

## USER

Ya

## ASSISTANT

Certainly, here is a Python script that aims to:

1. Load two pre-trained models (`finetuned\_bert\_with\_validation.pt` and `pytorch\_model.bin`)
2. Load the new training data from the provided CSV files
3. Train an ensemble model that uses the pre-trained models and the new data
4. Save the ensemble model as a `joblib` file

Before running this script, make sure you have uploaded the pre-trained model files (`finetuned\_bert\_with\_validation.pt` and `pytorch\_model.bin`) and the CSV files to your Google Colab workspace.

```
```python
```

```
import os
```

```
import joblib

import pandas as pd

import torch

from torch.utils.data import DataLoader, TensorDataset

from transformers import BertTokenizer, BertForSequenceClassification, AdamW

from sklearn.ensemble import VotingRegressor

from sklearn.metrics import mean_squared_error


# Load pre-trained models

finetuned_bert_path = '/content/drive/MyDrive/000/finetuned_bert_with_validation.pt'

pytorch_model_path = '/content/drive/MyDrive/000/pytorch_model.bin'

config_path = '/content/drive/MyDrive/000/config.json'


finetuned_bert_model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
num_labels=2)

pytorch_model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
num_labels=2)


# Load the pre-trained model weights

finetuned_bert_model.load_state_dict(torch.load(finetuned_bert_path))

pytorch_model.load_state_dict(torch.load(pytorch_model_path))


# Load new training data

train_data_path = '/content/california_housing_train.csv'

test_data_path = '/content/california_housing_test.csv'
```

```
train_df = pd.read_csv(train_data_path)
```

```
test_df = pd.read_csv(test_data_path)
```

```
# Assuming "median_house_value" is the target column
```

```
X_train = train_df.drop('median_house_value', axis=1)
```

```
y_train = train_df['median_house_value']
```

```
X_test = test_df.drop('median_house_value', axis=1)
```

```
y_test = test_df['median_house_value']
```

```
# Tokenize and prepare DataLoader for training new model
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
```

```
tokens_train    = tokenizer(list(X_train['ocean_proximity']), padding=True, truncation=True,  
return_tensors='pt')
```

```
tokens_test     = tokenizer(list(X_test['ocean_proximity']), padding=True, truncation=True,  
return_tensors='pt')
```

```
input_ids_train = tokens_train['input_ids']
```

```
attention_mask_train = tokens_train['attention_mask']
```

```
input_ids_test = tokens_test['input_ids']
```

```
attention_mask_test = tokens_test['attention_mask']
```

```
# Create a DataLoader
```

```
train_dataset = TensorDataset(input_ids_train, attention_mask_train, torch.tensor(y_train.values))
```

```
test_dataset = TensorDataset(input_ids_test, attention_mask_test, torch.tensor(y_test.values))
```

```
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
```

```
test_loader = DataLoader(test_dataset, batch_size=32)
```

```
# Initialize and train a new model
```

```
new_model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=1)
```

```
optimizer = AdamW(new_model.parameters(), lr=1e-5)
```

```
for epoch in range(3): # Perform 3 epochs
```

```
    for batch in train_loader:
```

```
        batch_input_ids, batch_attention_mask, batch_labels = batch
```

```
        optimizer.zero_grad()
```

```
            outputs = new_model(input_ids=batch_input_ids, attention_mask=batch_attention_mask,
```

```
labels=batch_labels.float())
```

```
            loss = outputs.loss
```

```
            loss.backward()
```

```
            optimizer.step()
```

```
# Ensemble model using VotingRegressor
```

```
ensemble_model = VotingRegressor([('finetuned_bert', finetuned_bert_model),
```

```
                                  ('pytorch_model', pytorch_model),
```

```
                                  ('new_model', new_model)])
```

```
ensemble_model.fit(X_train, y_train)
```

```
# Save the ensemble model
```

```
joblib.dump(ensemble_model, '/content/ensemble_model.joblib')
```

```
# Test ensemble model

ensemble_predictions = ensemble_model.predict(X_test)

mse = mean_squared_error(y_test, ensemble_predictions)

print(f'Mean Squared Error of Ensemble Model: {mse}')

...

```

Please note:

- This script assumes the column `ocean_proximity` is the feature to be used for the text-based model. You can replace it as needed.
- The script is designed to run in a Google Colab environment, so paths are set accordingly.
- You might need to install the `transformers` package by running `!pip install transformers`.